

# PPHA 30546 Machine Learning Problem Set 1

Richard Campo in collaboration with Julia Klauss

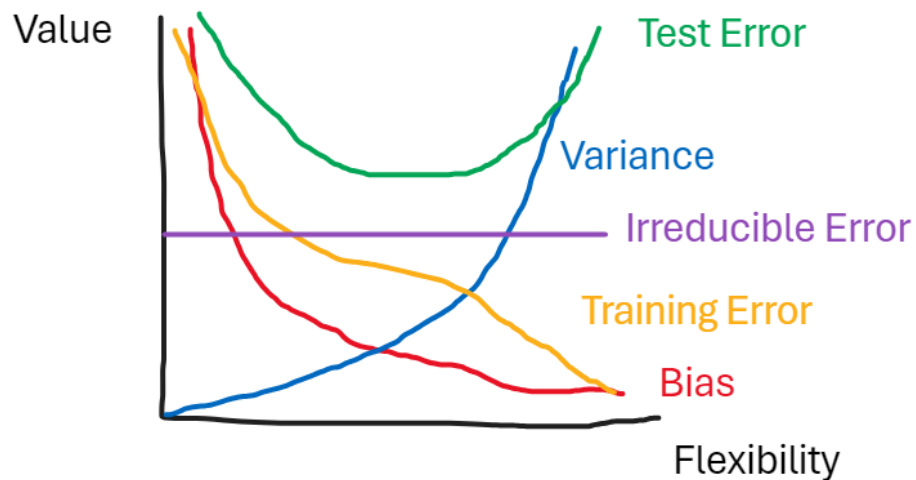
```
In [1]: from IPython.display import Image
import pandas as pd
import os
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
from IPython.display import display
```

## Chapter 2: Question 3

(a)

```
In [4]: # https://stackoverflow.com/a/11855133
Image(filename="Machine Learning Pset 1 Curves.png")
```

Out[4]:



(b)

- Bias decreases as flexibility increases because the model can overfit the data. Bias tends to decrease faster at lower levels of flexibility because adding just a little curvature can

allow the model to fit the data much better, but once you have lots of flexibility, adding more won't do much.

- Variance increases as flexibility increases because the model is more sensitive to new training data. Variance tends to increase at an increasing rate because very flexible models will change more than less flexible models when provided new training data due to overfitting.
- Training error decreases as flexibility increases because the model can fit the data more closely. Eventually, the most flexible model can correctly predict every point in the data by overfitting.
- Test error first decreases and then increases as flexibility increases. When flexibility is low, adding flexibility allows the model to more closely fit the training data and accurately predict the test data. However, when flexibility is high, adding flexibility causes overfitting, which lowers training error but increases test error.
- Irreducible error is a constant =  $\text{Var}(\epsilon)$

## Chapter 2: Question 5

The advantage of a very flexible approach is that the model can more closely fit the training data and more accurately predict the test data. However, the disadvantage of a very flexible approach is that the model can overfit the training data, meaning the model will make very accurate predictions on the training data, but not generalize to the test data.

A more flexible approach is preferred when the true underlying relationship in the data is very nonlinear because a line can't fit the data well. On the other hand, a less flexible approach is preferred when the true underlying relationship in the data is linear or approximately linear because flexible methods will overfit the data.

## Chapter 2: Question 10

(a)

```
In [2]: PATH = r"C:\Users\RichardCampo\Documents\GitHub\Machine-Learning\Boston"
df_boston = pd.read_csv(os.path.join(PATH, "Boston.csv"))
```

(b)

```
In [3]: print(df_boston.shape)
df_boston.head()
```

(506, 14)


| Out[3]: |         | CRIM | ZN   | INDUS | CHAS  | NOX   | RM   | AGE    | DIS | RAD   | TAX  | PTRATIO | B |
|---------|---------|------|------|-------|-------|-------|------|--------|-----|-------|------|---------|---|
| 0       | 0.00632 | 18.0 | 2.31 | 0.0   | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90  |   |
| 1       | 0.02731 | 0.0  | 7.07 | 0.0   | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90  |   |
| 2       | 0.02729 | 0.0  | 7.07 | 0.0   | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83  |   |
| 3       | 0.03237 | 0.0  | 2.18 | 0.0   | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63  |   |
| 4       | 0.06905 | 0.0  | 2.18 | 0.0   | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90  |   |

There are 506 rows and 14 columns in the dataset. Rows represent census tracts in the Boston area and columns represent variables describing each census tract. Additional information about the variables is available in the data description text file.

(c)

```
In [4]: # Plot first 5 variables to make it easier to read
sns.pairplot(df_boston.iloc[:, :5])
```

Out[4]: <seaborn.axisgrid.PairGrid at 0x1c12913e2f0>



Looking at the scatterplots, there is a positive correlations between nitric oxide concentration and the number of acres zoned for industry in a census district, which is unsurprising. It also looks like there is a negative association between the number of acres zoned for residential buildings and the amount of nitric oxide and amount of acres zoned for industry. This also makes sense. Lastly, per capita crime appears to be positively correlated with the concentration of nitric oxides. We could make the same plots using the other variables in the data set, but these have been excluded to save space.

(d)

```
In [5]: df_boston.corr()["CRIM"].sort_values(ascending=False)
```

```
Out[5]: CRIM      1.000000
        RAD      0.622029
        TAX      0.579564
        LSTAT    0.452220
        NOX      0.417521
        INDUS    0.404471
        AGE      0.350784
        PTRATIO   0.288250
        CHAS     -0.055295
        ZN       -0.199458
        RM       -0.219940
        B        -0.377365
        DIS      -0.377904
        MDEV     -0.385832
        Name: CRIM, dtype: float64
```


Most of the predictors are at least weakly associated with crime. The strongest correlation is between per capita crime and access to radial highways, maybe because areas close to highways are poorer or maybe because criminals use highways to move illegal drugs. Property tax rate is also positively associated with crime, perhaps because property taxes are higher closer to the city center where crime is more prevalent. Lastly, LSTAT, the percentage of people in the census district who are "lower status" (presumably in terms of income), is also positively correlated with per capita crime, which is what we expect.

(e)

```
In [42]: print(df_boston[["CRIM", "TAX", "PTRATIO"]].describe())
        df_boston[["CRIM", "TAX", "PTRATIO"]].hist()
```

|       | CRIM       | TAX        | PTRATIO    |
|-------|------------|------------|------------|
| count | 506.000000 | 506.000000 | 506.000000 |
| mean  | 3.593761   | 408.237154 | 18.455534  |
| std   | 8.596783   | 168.537116 | 2.164946   |
| min   | 0.006320   | 187.000000 | 12.600000  |
| 25%   | 0.082045   | 279.000000 | 17.400000  |
| 50%   | 0.256510   | 330.000000 | 19.050000  |
| 75%   | 3.647423   | 666.000000 | 20.200000  |
| max   | 88.976200  | 711.000000 | 22.000000  |

```
Out[42]: array([[<Axes: title={'center': 'CRIM'}>,
                  <Axes: title={'center': 'TAX'}>],
                [<Axes: title={'center': 'PTRATIO'}>, <Axes: >]], dtype=object)
```



Most census tracts have low crime per capita since the mean is only about 3.6. However, a small number of census tracts have high per capita crime rates, up to a maximum of almost 89.

Census tracts in Boston have full-value property tax rates per \$10,000 that vary from 187 to 711, but there is a large valley between about 450 and 650. A large group of census tracts have high property tax rates, possibly because they are close to the city center.

Lastly, pupil-teacher ratios vary from 12.6 to 22 pupils per teacher, with a large number of census districts having 20 to 21 pupils per teacher. These higher ratio census districts may also be closer to the city center where schools tend to have less funding per student.

(f)

```
In [6]: borders = len(df_boston[df_boston["CHAS"] == 1])

print(borders)
print(borders/len(df_boston) * 100, "%")
```

```
35
6.91699604743083 %
```

35 of the 506 census tracts border the Charles River, or about 7% of census tracts in the dataset.

(g)

```
In [7]: df_boston["PTRATIO"].median()
```

```
Out[7]: 19.05
```

The median pupil-teacher ratio of census tracts in the dataset is 19.05 pupils per teacher.

(h)

```
In [11]: df_boston.loc[[df_boston["MDEV"].idxmin()]]
```

```
Out[11]:
```

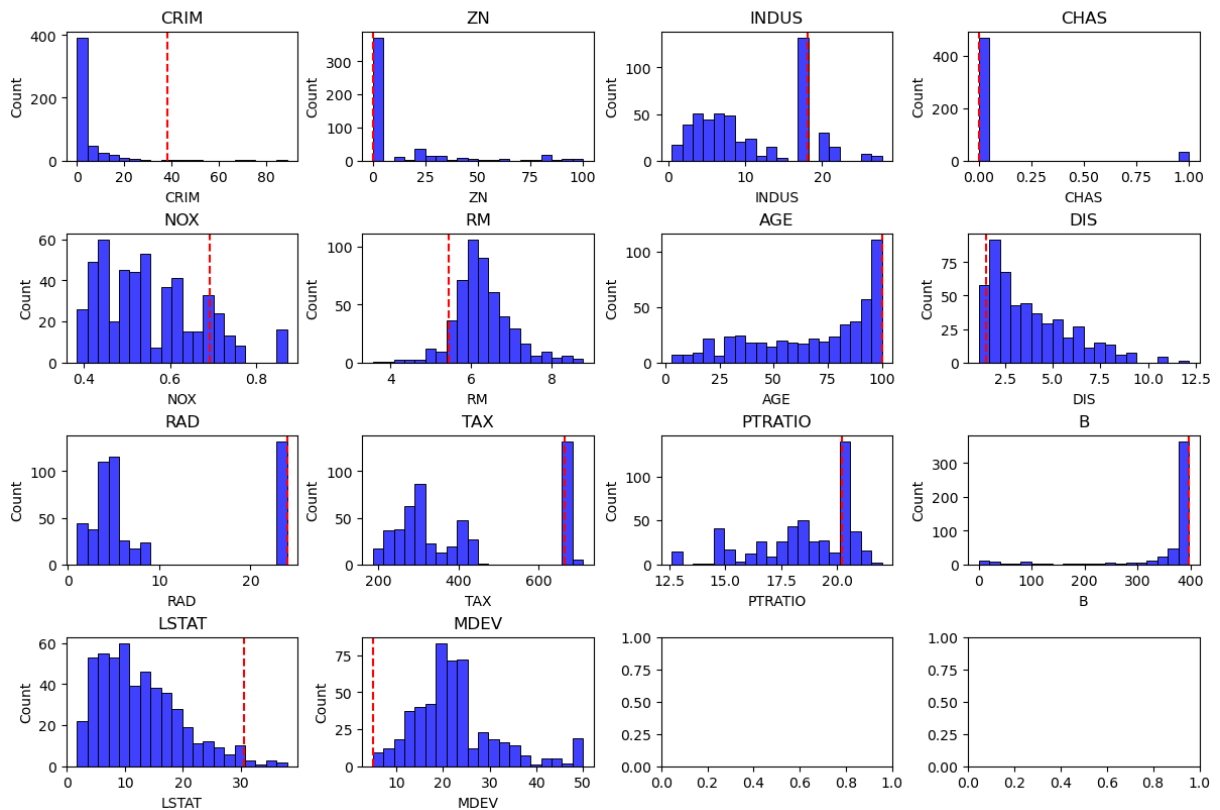
|     | CRIM    | ZN  | INDUS | CHAS | NOX   | RM    | AGE   | DIS    | RAD  | TAX   | PTRATIO | B     |
|-----|---------|-----|-------|------|-------|-------|-------|--------|------|-------|---------|-------|
| 398 | 38.3518 | 0.0 | 18.1  | 0.0  | 0.693 | 5.453 | 100.0 | 1.4896 | 24.0 | 666.0 | 20.2    | 396.9 |

```
In [46]: min_row_index = df_boston["MDEV"].idxmin()
min_row = df_boston.loc[[min_row_index]]

# https://stackoverflow.com/a/53649492
fig, axs = plt.subplots(4, 4, figsize=(12, 8), constrained_layout=True)
fig.suptitle("Histograms with Min MDEV Tract Highlighted", fontsize=16, y=1.05)

# Plot a histogram showing distribution for each variable with red line
# showing the value for the census tract with lowest MDEV.
for ax, column in zip(axs.flatten(), df_boston.columns):
    sns.histplot(df_boston[column], bins=20, color='blue', ax=ax)
    ax.axvline(x=min_row[column].values[0], color='red', linestyle='dashed')
    ax.set_title(column)
```

Histograms with Min MDEV Tract Highlighted



The census tract with the lowest median value of owner-occupied homes is located in row 398 and has a median home value of 5,000. *This is much lower than the median value of 21.2 thousand.* This census tract also has high crime per capita, an above average amount of industrially zoned space, above average nitric oxide concentration, low average number of rooms per building, very old buildings, low distance to the city employment centers, high accessibility to radial highways, high property taxes, a high pupil-teacher ratio, and a high proportion of people with low incomes.

This helps show that median home value is correlated with many of the other variables in our data set, so we want to use multiple regression to isolate the partial effects of each variable.

(i)

```
In [47]: print(len(df_boston.loc[df_boston["RM"] > 7]))
print(len(df_boston.loc[df_boston["RM"] > 8]))
```

64  
13

There are 64 census tracts where the average number of rooms per home is greater than 7, and there are 13 census tracts where the average number of rooms per home is more than 8.

```
In [48]: more_than_8 = df_boston.loc[df_boston["RM"] > 8].mean()
avg_values = df_boston.mean()
```



```
comparison = pd.concat([more_than_8, avg_values], axis = 1)
comparison.columns = ["more_than_8", "avg_values"]

print(comparison)
```

|         | more_than_8 | avg_values |
|---------|-------------|------------|
| CRIM    | 0.718795    | 3.593761   |
| ZN      | 13.615385   | 11.363636  |
| INDUS   | 7.078462    | 11.136779  |
| CHAS    | 0.153846    | 0.069170   |
| NOX     | 0.539238    | 0.554695   |
| RM      | 8.348538    | 6.284634   |
| AGE     | 71.538462   | 68.574901  |
| DIS     | 3.430192    | 3.795043   |
| RAD     | 7.461538    | 9.549407   |
| TAX     | 325.076923  | 408.237154 |
| PTRATIO | 16.361538   | 18.455534  |
| B       | 385.210769  | 356.674032 |
| LSTAT   | 4.310000    | 12.653063  |
| MDEV    | 44.200000   | 22.532806  |

Census tracts with more than eight rooms per dwelling tend to have low crime, have less industrially zoned land, are more likely to border the Charles River, have a smaller share of low income residents, and have high median home values. These are likely the fancier, wealthier areas of Boston where people can afford large, expensive homes.

## Chapter 3: Question 3

(a)

The "true" regression model is:

$$salary = \beta_0 + \beta_1 GPA + \beta_2 IQ + \beta_3 level + \beta_4 GPA \times IQ + \beta_5 GPA \times level + \epsilon$$

We estimate the model:

$$\widehat{salary} = 50 + 20(GPA) + 0.07(IQ) + 35(level) + 0.01(GPA \times IQ) - 10(GPA \times level)$$

The answer is (ii). For a fixed value of IQ and GPA, college graduates earn more on average than high school graduates. This is because  $\beta_3$ , the coefficient of level is 35, meaning a college graduate can expect a \$35,000 higher starting salary than a high school graduate on average, all else constant. The estimated coefficient of -10 for the interaction of GPA and level indicates that for college graduates, the slope of GPA is more shallow, but we are interested in the average difference between levels for this question, not GPA.

(b)

Based on our model, we predict that the starting salary of a college graduate with an IQ of 110 and a GPA of 4.0 will be

$$50 + 20(4.0) + 0.07(110) + 35(1) + 0.01(4.0)(110) - 10(4.0)(1) = 137.1$$

or \$137.1 thousand.

(c)

False. Just because the coefficient for the interaction between GPA and IQ is small does not mean there is little evidence of an interaction effect. The formula for a t-test to check if the coefficient is statistically significant is:

$$t = \frac{\hat{\beta}_j - \beta_j}{SE(\hat{\beta}_j)}$$

This means that even if the estimated coefficient is small, as long as the standard error is also very small, the coefficient could be statistically significant.

## Chapter 3: Question 15

(a)

```
In [8]: predictors = list(df_boston.columns)
predictors.remove("CRIM")

univar_coefs = {} # for part (c)

for predictor in predictors:
    X = df_boston[predictor]
    X = sm.add_constant(X) # add intercept to regression model
    y = df_boston["CRIM"]

    model = sm.OLS(y, X).fit()

    # Extract coefficient variable names and values for part (c)
    univar_coefs[predictor] = model.params[predictor]

    print(f"Regression results for {predictor}:")
    display(model.summary())
    print("\n")
```

Regression results for ZN:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.040    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.038    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 20.88    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 6.15e-06 |
| <b>Time:</b>             | 14:18:54         | <b>Log-Likelihood:</b>     | -1795.8  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3596.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3604.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | coef    | std err | t      | P> t  | [0.025 | 0.975] |
|--------------|---------|---------|--------|-------|--------|--------|
| <b>const</b> | 4.4292  | 0.417   | 10.620 | 0.000 | 3.610  | 5.249  |
| <b>ZN</b>    | -0.0735 | 0.016   | -4.570 | 0.000 | -0.105 | -0.042 |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 568.366 | <b>Durbin-Watson:</b>    | 0.862     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 32952.356 |
| <b>Skew:</b>          | 5.270   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 41.103  | <b>Cond. No.</b>         | 28.8      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for INDUS:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.164    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.162    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 98.58    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 2.44e-21 |
| <b>Time:</b>             | 14:18:54         | <b>Log-Likelihood:</b>     | -1760.9  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3526.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3534.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                       | coef    | std err                  | t         | P> t  | [0.025 | 0.975] |
|-----------------------|---------|--------------------------|-----------|-------|--------|--------|
| <b>const</b>          | -2.0509 | 0.668                    | -3.072    | 0.002 | -3.362 | -0.739 |
| <b>INDUS</b>          | 0.5068  | 0.051                    | 9.929     | 0.000 | 0.407  | 0.607  |
| <b>Omnibus:</b>       | 585.528 | <b>Durbin-Watson:</b>    | 0.990     |       |        |        |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 41469.710 |       |        |        |
| <b>Skew:</b>          | 5.456   | <b>Prob(JB):</b>         | 0.00      |       |        |        |
| <b>Kurtosis:</b>      | 45.987  | <b>Cond. No.</b>         | 25.1      |       |        |        |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for CHAS:

# OLS Regression Results

|                          |                  |                            |         |
|--------------------------|------------------|----------------------------|---------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.003   |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.001   |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 1.546   |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 0.214   |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1805.3 |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3615.   |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3623.   |
| <b>Df Model:</b>         | 1                |                            |         |
| <b>Covariance Type:</b>  | nonrobust        |                            |         |

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | 3.7232      | 0.396          | 9.404    | 0.000           | 2.945         | 4.501         |
| <b>CHAS</b>  | -1.8715     | 1.505          | -1.243   | 0.214           | -4.829        | 1.086         |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 562.698 | <b>Durbin-Watson:</b>    | 0.822     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 30864.755 |
| <b>Skew:</b>          | 5.205   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 39.818  | <b>Cond. No.</b>         | 3.96      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for NOX:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.174    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.173    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 106.4    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 9.16e-23 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1757.6  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3519.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3528.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                       | coef     | std err                  | t         | P> t  | [0.025  | 0.975]  |
|-----------------------|----------|--------------------------|-----------|-------|---------|---------|
| <b>const</b>          | -13.5881 | 1.702                    | -7.986    | 0.000 | -16.931 | -10.245 |
| <b>NOX</b>            | 30.9753  | 3.003                    | 10.315    | 0.000 | 25.076  | 36.875  |
| <b>Omnibus:</b>       | 591.496  | <b>Durbin-Watson:</b>    | 0.994     |       |         |         |
| <b>Prob(Omnibus):</b> | 0.000    | <b>Jarque-Bera (JB):</b> | 42994.381 |       |         |         |
| <b>Skew:</b>          | 5.544    | <b>Prob(JB):</b>         | 0.00      |       |         |         |
| <b>Kurtosis:</b>      | 46.776   | <b>Cond. No.</b>         | 11.3      |       |         |         |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for RM:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.048    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.046    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 25.62    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 5.84e-07 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1793.5  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3591.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3600.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | 20.5060     | 3.362          | 6.099    | 0.000           | 13.901        | 27.111        |
| <b>RM</b>    | -2.6910     | 0.532          | -5.062   | 0.000           | -3.736        | -1.646        |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 576.890 | <b>Durbin-Watson:</b>    | 0.883     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 36966.825 |
| <b>Skew:</b>          | 5.361   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 43.477  | <b>Cond. No.</b>         | 58.4      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for AGE:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.123    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.121    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 70.72    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 4.26e-16 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1772.9  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3550.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3558.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | -3.7527     | 0.944          | -3.974   | 0.000           | -5.608        | -1.898        |
| <b>AGE</b>   | 0.1071      | 0.013          | 8.409    | 0.000           | 0.082         | 0.132         |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 575.090 | <b>Durbin-Watson:</b>    | 0.960     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 36851.412 |
| <b>Skew:</b>          | 5.331   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 43.426  | <b>Cond. No.</b>         | 195.      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for DIS:



# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.143    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.141    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 83.97    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.27e-18 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1767.1  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3538.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3547.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | 9.4489      | 0.731          | 12.934   | 0.000           | 8.014         | 10.884        |
| <b>DIS</b>   | -1.5428     | 0.168          | -9.163   | 0.000           | -1.874        | -1.212        |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 577.090 | <b>Durbin-Watson:</b>    | 0.957     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 37542.100 |
| <b>Skew:</b>          | 5.357   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 43.815  | <b>Cond. No.</b>         | 9.32      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for RAD:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.387    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.386    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 318.1    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.62e-55 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1682.3  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3369.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3377.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | -2.2709     | 0.445          | -5.105   | 0.000           | -3.145        | -1.397        |
| <b>RAD</b>   | 0.6141      | 0.034          | 17.835   | 0.000           | 0.546         | 0.682         |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 654.232 | <b>Durbin-Watson:</b>    | 1.336     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 74327.568 |
| <b>Skew:</b>          | 6.441   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 60.961  | <b>Cond. No.</b>         | 19.2      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for TAX:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.336    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.335    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 254.9    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 9.76e-47 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1702.5  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3409.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3418.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | coef    | std err | t       | P> t  | [0.025  | 0.975] |
|--------------|---------|---------|---------|-------|---------|--------|
| <b>const</b> | -8.4748 | 0.818   | -10.365 | 0.000 | -10.081 | -6.868 |
| <b>TAX</b>   | 0.0296  | 0.002   | 15.966  | 0.000 | 0.026   | 0.033  |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 634.003 | <b>Durbin-Watson:</b>    | 1.252     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 63141.063 |
| <b>Skew:</b>          | 6.134   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 56.332  | <b>Cond. No.</b>         | 1.16e+03  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.16e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for PTRATIO:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.083    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.081    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 45.67    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 3.88e-11 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1784.1  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3572.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3581.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|----------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b>   | -17.5307    | 3.147          | -5.570   | 0.000           | -23.714       | -11.347       |
| <b>PTRATIO</b> | 1.1446      | 0.169          | 6.758    | 0.000           | 0.812         | 1.477         |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 568.808 | <b>Durbin-Watson:</b>    | 0.909     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 34373.378 |
| <b>Skew:</b>          | 5.256   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 41.985  | <b>Cond. No.</b>         | 160.      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for B:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.142    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.141    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 83.69    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.43e-18 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1767.2  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3538.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3547.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                       | <b>coef</b> | <b>std err</b>           | <b>t</b>  | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|-----------------------|-------------|--------------------------|-----------|-----------------|---------------|---------------|
| <b>const</b>          | 16.2680     | 1.430                    | 11.376    | 0.000           | 13.458        | 19.078        |
| <b>B</b>              | -0.0355     | 0.004                    | -9.148    | 0.000           | -0.043        | -0.028        |
| <b>Omnibus:</b>       | 591.626     | <b>Durbin-Watson:</b>    | 1.001     |                 |               |               |
| <b>Prob(Omnibus):</b> | 0.000       | <b>Jarque-Bera (JB):</b> | 43282.465 |                 |               |               |
| <b>Skew:</b>          | 5.543       | <b>Prob(JB):</b>         | 0.00      |                 |               |               |
| <b>Kurtosis:</b>      | 46.932      | <b>Cond. No.</b>         | 1.49e+03  |                 |               |               |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.49e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for LSTAT:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.205    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.203    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 129.6    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 7.12e-27 |
| <b>Time:</b>             | 14:18:55         | <b>Log-Likelihood:</b>     | -1748.2  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3500.    |
| <b>Df Residuals:</b>     | 504              | <b>BIC:</b>                | 3509.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | coef    | std err | t      | P> t  | [0.025 | 0.975] |
|--------------|---------|---------|--------|-------|--------|--------|
| <b>const</b> | -3.2946 | 0.695   | -4.742 | 0.000 | -4.660 | -1.930 |
| <b>LSTAT</b> | 0.5444  | 0.048   | 11.383 | 0.000 | 0.450  | 0.638  |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 600.766 | <b>Durbin-Watson:</b>    | 1.184     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 49637.173 |
| <b>Skew:</b>          | 5.638   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 50.193  | <b>Cond. No.</b>         | 29.7      |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for MDEV:

# OLS Regression Results

|                   |         |                  |                     |       |               |
|-------------------|---------|------------------|---------------------|-------|---------------|
| Dep. Variable:    |         | CRIM             | R-squared:          |       | 0.149         |
| Model:            |         | OLS              | Adj. R-squared:     |       | 0.147         |
| Method:           |         | Least Squares    | F-statistic:        |       | 88.15         |
| Date:             |         | Wed, 17 Jan 2024 | Prob (F-statistic): |       | 2.08e-19      |
| Time:             |         | 14:18:55         | Log-Likelihood:     |       | -1765.3       |
| No. Observations: |         | 506              | AIC:                |       | 3535.         |
| Df Residuals:     |         | 504              | BIC:                |       | 3543.         |
| Df Model:         |         | 1                |                     |       |               |
| Covariance Type:  |         | nonrobust        |                     |       |               |
|                   | coef    | std err          | t                   | P> t  | [0.025 0.975] |
| const             | 11.7202 | 0.935            | 12.539              | 0.000 | 9.884 13.557  |
| MDEV              | -0.3606 | 0.038            | -9.389              | 0.000 | -0.436 -0.285 |
| Omnibus:          |         | 559.282          | Durbin-Watson:      |       | 1.000         |
| Prob(Omnibus):    |         | 0.000            | Jarque-Bera (JB):   |       | 32809.507     |
| Skew:             |         | 5.114            | Prob(JB):           |       | 0.00          |
| Kurtosis:         |         | 41.099           | Cond. No.           |       | 64.5          |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Almost every predictor on its own is significantly correlated with the per capita crime rate, as we can see from the p-values that are substantially lower than 0.05. The only exception is whether the census tract borders the Charles River, which is weakly correlated with lower per capita crime, but the estimated coefficient is not statistically significant at the 5% level.

```
In [50]: for predictor in predictors:
          X = df_boston[predictor]
          X = sm.add_constant(X)
          y = df_boston["CRIM"]

          model = sm.OLS(y, X).fit()

          # Create Figure and Axes objects
          fig, ax = plt.subplots(figsize=(6, 4))

          # Plot the data points
```

```

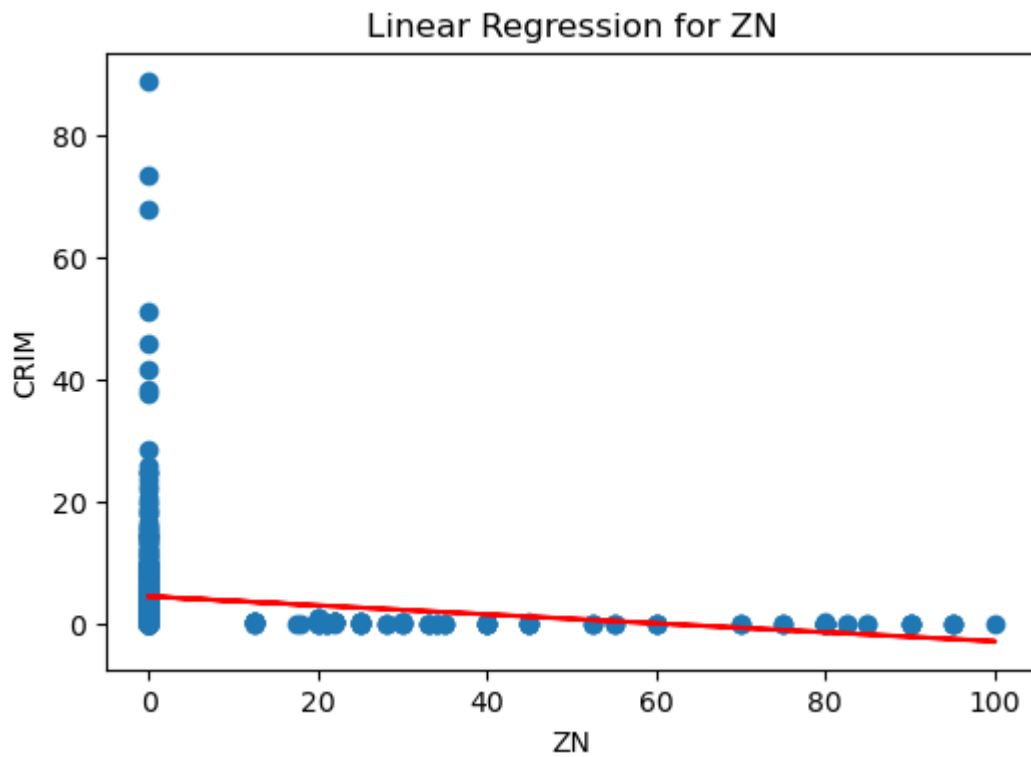
scatter = ax.scatter(
    X.iloc[:, 1], # exclude intercept column
    y
)

# Plot the regression line
line = ax.plot(
    X.iloc[:, 1], # exclude intercept column
    model.predict(X),
    color='red'
)

ax.set_title(f"Linear Regression for {predictor}")
ax.set_xlabel(predictor)
ax.set_ylabel("CRIM")

# Display the plot
plt.show()

```







A scatter plot showing the relationship between CRIM (Crime Rate) on the y-axis and INDUS (Non-retail business acres per city block) on the x-axis. The y-axis ranges from 0 to 80, and the x-axis ranges from 0 to 25. A red regression line is drawn through the data points, indicating a positive correlation. The data points are mostly clustered at low INDUS values (below 10) with CRIM values near 0. There is a significant outlier at INDUS ≈ 18.5, where CRIM values range from approximately 0 to 90.

A scatter plot showing the relationship between CRIM (Y-axis) and CHAS (X-axis). The Y-axis ranges from 0 to 80, and the X-axis ranges from 0.0 to 1.0. The data points are blue dots, and a red line represents the linear regression. The plot shows a high concentration of data points at CHAS=0.0, with CRIM values ranging from approximately 0 to 90. A small cluster of points is visible at CHAS=1.0, with CRIM values around 0 to 10. The red regression line indicates a very slight negative correlation between CRIM and CHAS.


Linear Regression for NOX




Linear Regression for RM




Linear Regression for AGE




Linear Regression for DIS



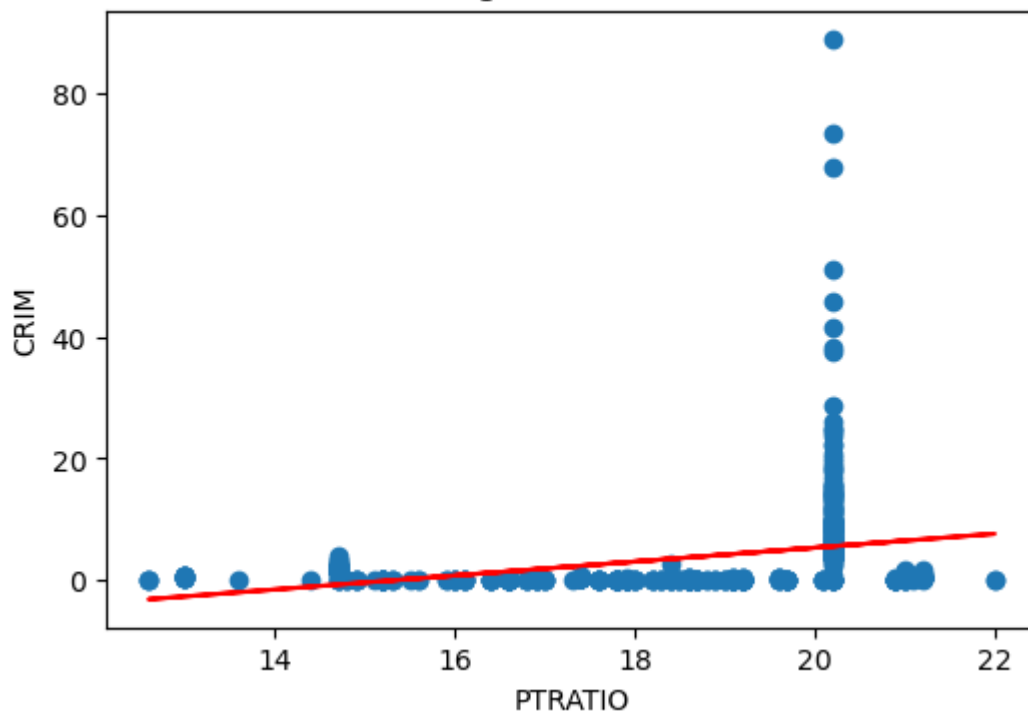
Linear Regression for RAD




Linear Regression for TAX

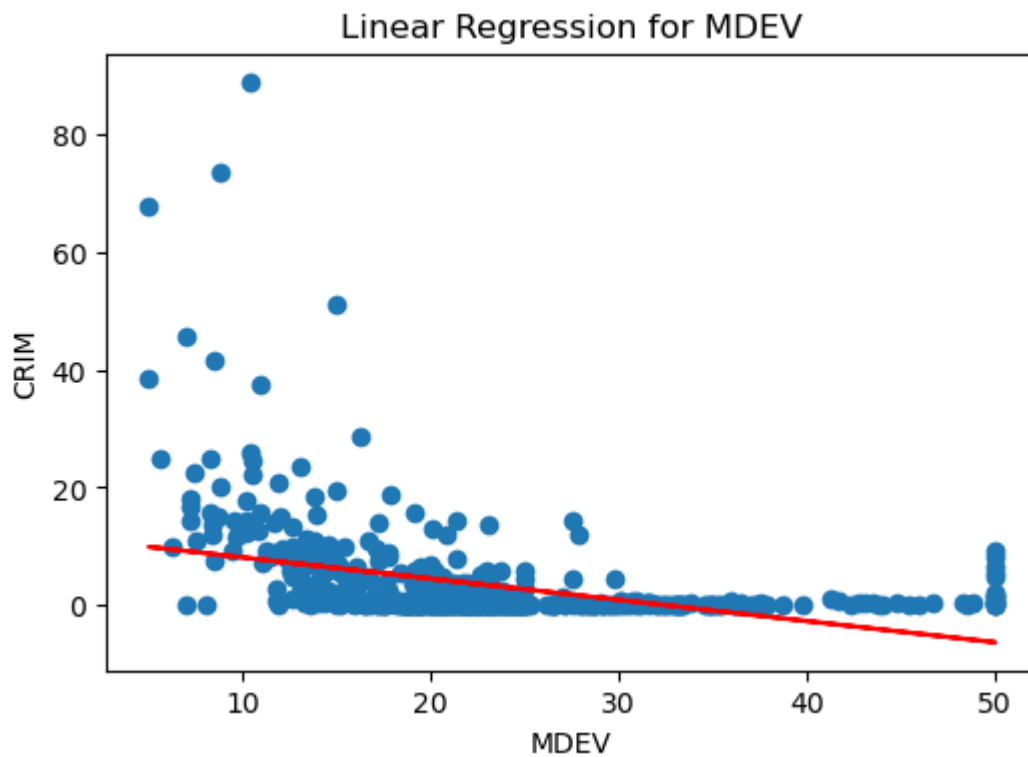



Linear Regression for PTRATIO



Linear Regression for B





We can see that there are some correlations for each of the predictors except "CHAS," the dummy for whether the census tract borders the Charles River. There is a weak negative correlation between "CHAS" and "CRIM," but it looks like most census tracts have relatively low crime and do not border the Charles River. The tracts with high crime per capita do not border the Charles River, however.

We can also see that many of the relationships between predictors and CRIM look nonlinear, which suggests we may want to use a polynomial regression to allow our line of best fit to curve to better fit the data.

**(b)**

```
In [9]: X = df_boston[predictors]
X = sm.add_constant(X)
y = df_boston["CRIM"]

model = sm.OLS(y, X).fit()
model.summary()
```

Out[9]:

## OLS Regression Results

|                   |          |                  |        |                     |         |           |  |
|-------------------|----------|------------------|--------|---------------------|---------|-----------|--|
| Dep. Variable:    |          | CRIM             |        | R-squared:          |         | 0.448     |  |
| Model:            |          | OLS              |        | Adj. R-squared:     |         | 0.434     |  |
| Method:           |          | Least Squares    |        | F-statistic:        |         | 30.73     |  |
| Date:             |          | Wed, 17 Jan 2024 |        | Prob (F-statistic): |         | 2.04e-55  |  |
| Time:             |          | 14:23:37         |        | Log-Likelihood:     |         | -1655.7   |  |
| No. Observations: |          | 506              |        | AIC:                |         | 3339.     |  |
| Df Residuals:     |          | 492              |        | BIC:                |         | 3399.     |  |
| Df Model:         |          | 13               |        |                     |         |           |  |
| Covariance Type:  |          | nonrobust        |        |                     |         |           |  |
|                   | coef     | std err          | t      | P> t                | [0.025  | 0.975]    |  |
| const             | 17.4184  | 7.270            | 2.396  | 0.017               | 3.135   | 31.702    |  |
| ZN                | 0.0449   | 0.019            | 2.386  | 0.017               | 0.008   | 0.082     |  |
| INDUS             | -0.0616  | 0.084            | -0.735 | 0.463               | -0.226  | 0.103     |  |
| CHAS              | -0.7414  | 1.186            | -0.625 | 0.532               | -3.071  | 1.588     |  |
| NOX               | -10.6455 | 5.301            | -2.008 | 0.045               | -21.061 | -0.230    |  |
| RM                | 0.3811   | 0.616            | 0.619  | 0.536               | -0.829  | 1.591     |  |
| AGE               | 0.0020   | 0.018            | 0.112  | 0.911               | -0.033  | 0.037     |  |
| DIS               | -0.9950  | 0.283            | -3.514 | 0.000               | -1.551  | -0.439    |  |
| RAD               | 0.5888   | 0.088            | 6.656  | 0.000               | 0.415   | 0.763     |  |
| TAX               | -0.0037  | 0.005            | -0.723 | 0.470               | -0.014  | 0.006     |  |
| PTRATIO           | -0.2787  | 0.187            | -1.488 | 0.137               | -0.647  | 0.089     |  |
| B                 | -0.0069  | 0.004            | -1.857 | 0.064               | -0.014  | 0.000     |  |
| LSTAT             | 0.1213   | 0.076            | 1.594  | 0.112               | -0.028  | 0.271     |  |
| MDEV              | -0.1992  | 0.061            | -3.276 | 0.001               | -0.319  | -0.080    |  |
| Omnibus:          |          | 662.271          |        | Durbin-Watson:      |         | 1.515     |  |
| Prob(Omnibus):    |          | 0.000            |        | Jarque-Bera (JB):   |         | 82701.666 |  |
| Skew:             |          | 6.544            |        | Prob(JB):           |         | 0.00      |  |
| Kurtosis:         |          | 64.248           |        | Cond. No.           |         | 1.58e+04  |  |



Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.58e+04. This might indicate that there are strong multicollinearity or other numerical problems.

When we just look at the partial effects of each predictor, fewer are predictive of the crime rate per capita. Now, only amount of residential zoning, nitric oxide emissions, distance from employment centers, access to radial highways, and median home value are significant at the 5% significance level.

For these predictors, we can reject the null hypothesis that there is no predicted effect of the independent variable on the per capita crime rate in Boston, holding the other predictors constant. We have evidence in favor of the alternative hypothesis that these predictors do have an expected effect on Boston per capita crime rates, holding the other predictors constant.

For the predictors with p-values greater than 0.05: INDUS, CHAS, RM, AGE, TAX, PTRATIO, and LSTAT, we fail to reject the null hypothesis that these variables have no expected effect on crime rate per capita.

(c)

```
In [11]: multi_coefs = model.params.drop("const")

# https://stackoverflow.com/a/47561390
uni_vs_multi = pd.DataFrame(
    {"univariate coefficients": pd.Series(univar_coefs), # from part (a)
     "multivariate coefficients": multi_coefs}
)

uni_vs_multi
```

Out[11]:

|                | univariate coefficients | multivariate coefficients |
|----------------|-------------------------|---------------------------|
| <b>ZN</b>      | -0.073521               | 0.044919                  |
| <b>INDUS</b>   | 0.506847                | -0.061576                 |
| <b>CHAS</b>    | -1.871545               | -0.741435                 |
| <b>NOX</b>     | 30.975259               | -10.645500                |
| <b>RM</b>      | -2.691045               | 0.381070                  |
| <b>AGE</b>     | 0.107131                | 0.002011                  |
| <b>DIS</b>     | -1.542831               | -0.994992                 |
| <b>RAD</b>     | 0.614137                | 0.588838                  |
| <b>TAX</b>     | 0.029563                | -0.003746                 |
| <b>PTRATIO</b> | 1.144613                | -0.278731                 |
| <b>B</b>       | -0.035535               | -0.006855                 |
| <b>LSTAT</b>   | 0.544406                | 0.121269                  |
| <b>MDEV</b>    | -0.360647               | -0.199218                 |

```
In [12]: fig, ax = plt.subplots()

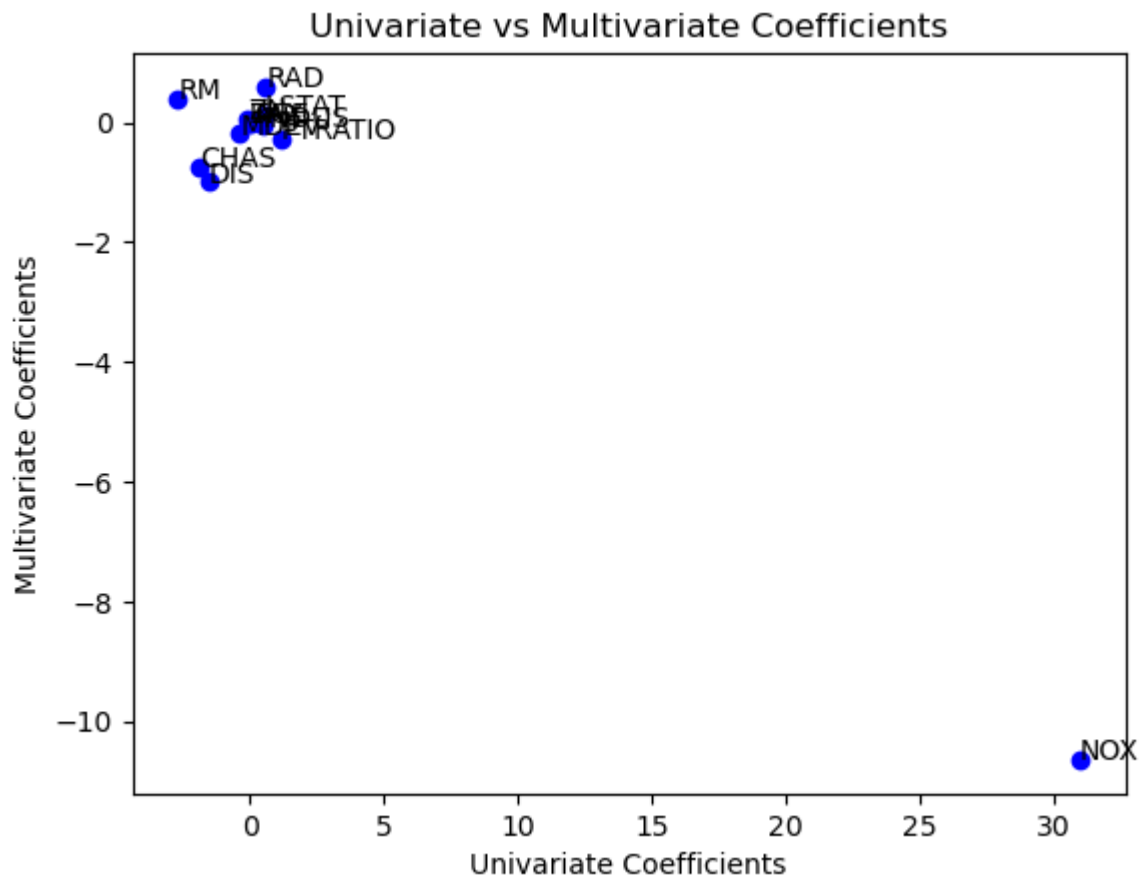
ax.scatter(
    uni_vs_multi['univariate coefficients'],
    uni_vs_multi['multivariate coefficients'],
    color='blue'
)

ax.set_title('Univariate vs Multivariate Coefficients')
ax.set_xlabel('Univariate Coefficients')
ax.set_ylabel('Multivariate Coefficients')

# Add labels for each point
for i, label in enumerate(uni_vs_multi.index):
    ax.annotate(
        label,
        # coordinates
        (uni_vs_multi['univariate coefficients'][i],
         uni_vs_multi['multivariate coefficients'][i])
    )

plt.show()

print(df_boston["NOX"].corr(df_boston["RAD"]))
```



0.6114405634855775

We can see from the scatterplot that most of the univariate coefficients were similar though more positive than their multivariate counterparts. The coefficient that changed the most was nitric oxide concentration, which was a strong positive predictor of a higher per capita crime rate in the univariate regression, but actually predicts a lower per capita crime rate when we isolate its partial effect.

This may be because areas with high nitric oxide concentrations tend to be more industrial and closer to the city center, so they tend to have more crime on average. However, high nitric oxide concentrations actually are associated with lower crime rates, all else constant. The difference comes from the fact that NOX is correlated with other variables that are positively associated with crime, such as access to radial highways, which were confounding the univariate model.

(d)

```
In [13]: for predictor in predictors:
          X = df_boston[predictor]
          X_squared = X**2
          X_cubed = X**3

          X_squared = X_squared.rename(f"{predictor}^2")
          X_cubed = X_cubed.rename(f"{predictor}^3")
```

```

X = pd.concat([X, X_squared, X_cubed], axis=1)
X = sm.add_constant(X)

y = df_boston["CRIM"]

model = sm.OLS(y, X).fit()

print(f"Regression results for {predictor}:")
display(model.summary())
print("\n")

```

Regression results for ZN:

#### OLS Regression Results

|                   |                  |                     |                                 |
|-------------------|------------------|---------------------|---------------------------------|
| Dep. Variable:    | CRIM             | R-squared:          | 0.058                           |
| Model:            | OLS              | Adj. R-squared:     | 0.052                           |
| Method:           | Least Squares    | F-statistic:        | 10.24                           |
| Date:             | Wed, 17 Jan 2024 | Prob (F-statistic): | 1.49e-06                        |
| Time:             | 14:48:39         | Log-Likelihood:     | -1791.1                         |
| No. Observations: | 506              | AIC:                | 3590.                           |
| Df Residuals:     | 502              | BIC:                | 3607.                           |
| Df Model:         | 3                |                     |                                 |
| Covariance Type:  | nonrobust        |                     |                                 |
|                   | coef             | std err             | t P> t  [0.025 0.975]           |
| const             | 4.8193           | 0.433               | 11.133 0.000 3.969 5.670        |
| ZN                | -0.3303          | 0.110               | -3.008 0.003 -0.546 -0.115      |
| ZN^2              | 0.0064           | 0.004               | 1.670 0.096 -0.001 0.014        |
| ZN^3              | -3.753e-05       | 3.14e-05            | -1.196 0.232 -9.92e-05 2.41e-05 |
| Omnibus:          | 570.003          | Durbin-Watson:      | 0.879                           |
| Prob(Omnibus):    | 0.000            | Jarque-Bera (JB):   | 33886.468                       |
| Skew:             | 5.285            | Prob(JB):           | 0.00                            |
| Kurtosis:         | 41.672           | Cond. No.           | 1.89e+05                        |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.89e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for INDUS:

OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.257    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.252    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 57.86    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 3.88e-32 |
| <b>Time:</b>             | 14:48:39         | <b>Log-Likelihood:</b>     | -1731.0  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3470.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3487.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                | coef    | std err | t      | P> t  | [0.025 | 0.975] |
|----------------|---------|---------|--------|-------|--------|--------|
| <b>const</b>   | 3.6410  | 1.576   | 2.310  | 0.021 | 0.545  | 6.737  |
| <b>INDUS</b>   | -1.9533 | 0.483   | -4.047 | 0.000 | -2.901 | -1.005 |
| <b>INDUS^2</b> | 0.2504  | 0.039   | 6.361  | 0.000 | 0.173  | 0.328  |
| <b>INDUS^3</b> | -0.0069 | 0.001   | -7.239 | 0.000 | -0.009 | -0.005 |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 611.416 | <b>Durbin-Watson:</b>    | 1.118     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 51547.097 |
| <b>Skew:</b>          | 5.815   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 51.059  | <b>Cond. No.</b>         | 2.47e+04  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.47e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for CHAS:

# OLS Regression Results

|                   |                  |                     |                       |                    |  |
|-------------------|------------------|---------------------|-----------------------|--------------------|--|
| Dep. Variable:    | CRIM             | R-squared:          | 0.003                 |                    |  |
| Model:            | OLS              | Adj. R-squared:     | -0.001                |                    |  |
| Method:           | Least Squares    | F-statistic:        | 0.7710                |                    |  |
| Date:             | Wed, 17 Jan 2024 | Prob (F-statistic): | 0.463                 |                    |  |
| Time:             | 14:48:39         | Log-Likelihood:     | -1805.3               |                    |  |
| No. Observations: | 506              | AIC:                | 3617.                 |                    |  |
| Df Residuals:     | 503              | BIC:                | 3629.                 |                    |  |
| Df Model:         | 2                |                     |                       |                    |  |
| Covariance Type:  | nonrobust        |                     |                       |                    |  |
|                   | coef             | std err             | t P> t  [0.025 0.975] |                    |  |
| const             | 3.7232           | 0.396               | 9.395 0.000           | 2.945 4.502        |  |
| CHAS              | 1.108e+14        | 2.71e+14            | 0.409 0.683           | -4.22e+14 6.43e+14 |  |
| CHAS^2            | -5.578e+13       | 1.36e+14            | -0.409 0.683          | -3.24e+14 2.12e+14 |  |
| CHAS^3            | -5.5e+13         | 1.35e+14            | -0.409 0.683          | -3.19e+14 2.09e+14 |  |
| Omnibus:          | 562.637          | Durbin-Watson:      | 0.822                 |                    |  |
| Prob(Omnibus):    | 0.000            | Jarque-Bera (JB):   | 30853.674             |                    |  |
| Skew:             | 5.204            | Prob(JB):           | 0.00                  |                    |  |
| Kurtosis:         | 39.811           | Cond. No.           | 5.51e+16              |                    |  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.7e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Regression results for NOX:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.292    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.288    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 69.14    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.94e-37 |
| <b>Time:</b>             | 14:48:39         | <b>Log-Likelihood:</b>     | -1718.6  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3445.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3462.    |
| <b>Df Model:</b>         | 3                |                            |          |

**Covariance Type:** nonrobust

|              | coef       | std err | t      | P> t  | [0.025    | 0.975]   |
|--------------|------------|---------|--------|-------|-----------|----------|
| <b>const</b> | 230.1421   | 33.734  | 6.822  | 0.000 | 163.864   | 296.420  |
| <b>NOX</b>   | -1264.1021 | 170.860 | -7.398 | 0.000 | -1599.791 | -928.414 |
| <b>NOX^2</b> | 2223.2265  | 280.659 | 7.921  | 0.000 | 1671.816  | 2774.637 |
| <b>NOX^3</b> | -1232.3894 | 149.687 | -8.233 | 0.000 | -1526.479 | -938.300 |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 612.604 | <b>Durbin-Watson:</b>    | 1.159     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 52872.508 |
| <b>Skew:</b>          | 5.824   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 51.705  | <b>Cond. No.</b>         | 1.36e+03  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.36e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for RM:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.068    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.063    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 12.29    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 9.06e-08 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1788.2  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3584.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3601.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | coef     | std err | t      | P> t  | [0.025   | 0.975]  |
|--------------|----------|---------|--------|-------|----------|---------|
| <b>const</b> | 111.9002 | 64.460  | 1.736  | 0.083 | -14.744  | 238.545 |
| <b>RM</b>    | -38.7040 | 31.284  | -1.237 | 0.217 | -100.167 | 22.759  |
| <b>RM^2</b>  | 4.4655   | 5.005   | 0.892  | 0.373 | -5.369   | 14.300  |
| <b>RM^3</b>  | -0.1694  | 0.264   | -0.643 | 0.521 | -0.687   | 0.348   |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 586.445 | <b>Durbin-Watson:</b>    | 0.919     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 40548.719 |
| <b>Skew:</b>          | 5.484   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 45.461  | <b>Cond. No.</b>         | 5.36e+04  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.36e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for AGE:



# OLS Regression Results

|                          |                  |                            |           |                 |               |               |
|--------------------------|------------------|----------------------------|-----------|-----------------|---------------|---------------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.172     |                 |               |               |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.167     |                 |               |               |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 34.86     |                 |               |               |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.76e-20  |                 |               |               |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1758.2   |                 |               |               |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3524.     |                 |               |               |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3541.     |                 |               |               |
| <b>Df Model:</b>         | 3                |                            |           |                 |               |               |
| <b>Covariance Type:</b>  | nonrobust        |                            |           |                 |               |               |
|                          | <b>coef</b>      | <b>std err</b>             | <b>t</b>  | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
| <b>const</b>             | -2.5592          | 2.771                      | -0.924    | 0.356           | -8.003        | 2.884         |
| <b>AGE</b>               | 0.2743           | 0.186                      | 1.471     | 0.142           | -0.092        | 0.641         |
| <b>AGE^2</b>             | -0.0072          | 0.004                      | -1.987    | 0.047           | -0.014        | -8.25e-05     |
| <b>AGE^3</b>             | 5.737e-05        | 2.11e-05                   | 2.719     | 0.007           | 1.59e-05      | 9.88e-05      |
| <b>Omnibus:</b>          | 577.859          | <b>Durbin-Watson:</b>      | 1.027     |                 |               |               |
| <b>Prob(Omnibus):</b>    | 0.000            | <b>Jarque-Bera (JB):</b>   | 39629.126 |                 |               |               |
| <b>Skew:</b>             | 5.342            | <b>Prob(JB):</b>           | 0.00      |                 |               |               |
| <b>Kurtosis:</b>         | 45.018           | <b>Cond. No.</b>           | 4.74e+06  |                 |               |               |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.74e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for DIS:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.276    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.272    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 63.74    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 6.20e-35 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1724.4  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3457.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3474.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|              | coef     | std err | t      | P> t  | [0.025  | 0.975]  |
|--------------|----------|---------|--------|-------|---------|---------|
| <b>const</b> | 29.9496  | 2.448   | 12.235 | 0.000 | 25.140  | 34.759  |
| <b>DIS</b>   | -15.5172 | 1.737   | -8.931 | 0.000 | -18.931 | -12.104 |
| <b>DIS^2</b> | 2.4479   | 0.347   | 7.061  | 0.000 | 1.767   | 3.129   |
| <b>DIS^3</b> | -0.1185  | 0.020   | -5.802 | 0.000 | -0.159  | -0.078  |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 577.986 | <b>Durbin-Watson:</b>    | 1.133     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 42441.952 |
| <b>Skew:</b>          | 5.310   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 46.592  | <b>Cond. No.</b>         | 2.10e+03  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.1e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for RAD:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.396    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.392    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 109.5    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 1.47e-54 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1678.7  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3365.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3382.    |
| <b>Df Model:</b>         | 3                |                            |          |

**Covariance Type:** nonrobust

|              | <b>coef</b> | <b>std err</b> | <b>t</b> | <b>P&gt; t </b> | <b>[0.025</b> | <b>0.975]</b> |
|--------------|-------------|----------------|----------|-----------------|---------------|---------------|
| <b>const</b> | -0.6050     | 2.057          | -0.294   | 0.769           | -4.645        | 3.435         |
| <b>RAD</b>   | 0.5122      | 1.047          | 0.489    | 0.625           | -1.545        | 2.569         |
| <b>RAD^2</b> | -0.0750     | 0.149          | -0.504   | 0.615           | -0.368        | 0.218         |
| <b>RAD^3</b> | 0.0032      | 0.005          | 0.699    | 0.485           | -0.006        | 0.012         |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 657.375 | <b>Durbin-Watson:</b>    | 1.349     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 76643.757 |
| <b>Skew:</b>          | 6.487   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 61.881  | <b>Cond. No.</b>         | 5.43e+04  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.43e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for TAX:

# OLS Regression Results

|                   |            |                  |                     |       |           |          |
|-------------------|------------|------------------|---------------------|-------|-----------|----------|
| Dep. Variable:    |            | CRIM             | R-squared:          |       | 0.365     |          |
| Model:            |            | OLS              | Adj. R-squared:     |       | 0.361     |          |
| Method:           |            | Least Squares    | F-statistic:        |       | 96.10     |          |
| Date:             |            | Wed, 17 Jan 2024 | Prob (F-statistic): |       | 3.69e-49  |          |
| Time:             |            | 14:48:40         | Log-Likelihood:     |       | -1691.3   |          |
| No. Observations: |            | 506              | AIC:                |       | 3391.     |          |
| Df Residuals:     |            | 502              | BIC:                |       | 3407.     |          |
| Df Model:         |            | 3                |                     |       |           |          |
| Covariance Type:  |            | nonrobust        |                     |       |           |          |
|                   | coef       | std err          | t                   | P> t  | [0.025    | 0.975]   |
| const             | 19.0705    | 11.827           | 1.612               | 0.107 | -4.166    | 42.307   |
| TAX               | -0.1524    | 0.096            | -1.589              | 0.113 | -0.341    | 0.036    |
| TAX^2             | 0.0004     | 0.000            | 1.476               | 0.141 | -0.000    | 0.001    |
| TAX^3             | -2.193e-07 | 1.89e-07         | -1.158              | 0.247 | -5.91e-07 | 1.53e-07 |
| Omnibus:          |            | 642.369          | Durbin-Watson:      |       | 1.292     |          |
| Prob(Omnibus):    |            | 0.000            | Jarque-Bera (JB):   |       | 68905.900 |          |
| Skew:             |            | 6.249            | Prob(JB):           |       | 0.00      |          |
| Kurtosis:         |            | 58.786           | Cond. No.           |       | 6.16e+09  |          |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.16e+09. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for PTRATIO:

# OLS Regression Results

|                   |                  |                     |                       |                  |
|-------------------|------------------|---------------------|-----------------------|------------------|
| Dep. Variable:    | CRIM             | R-squared:          | 0.112                 |                  |
| Model:            | OLS              | Adj. R-squared:     | 0.107                 |                  |
| Method:           | Least Squares    | F-statistic:        | 21.21                 |                  |
| Date:             | Wed, 17 Jan 2024 | Prob (F-statistic): | 5.99e-13              |                  |
| Time:             | 14:48:40         | Log-Likelihood:     | -1775.9               |                  |
| No. Observations: | 506              | AIC:                | 3560.                 |                  |
| Df Residuals:     | 502              | BIC:                | 3577.                 |                  |
| Df Model:         | 3                |                     |                       |                  |
| Covariance Type:  | nonrobust        |                     |                       |                  |
|                   | coef             | std err             | t P> t  [0.025 0.975] |                  |
| const             | 474.0255         | 156.823             | 3.023 0.003           | 165.915 782.135  |
| PTRATIO           | -81.8089         | 27.649              | -2.959 0.003          | -136.131 -27.487 |
| PTRATIO^2         | 4.6039           | 1.609               | 2.862 0.004           | 1.444 7.764      |
| PTRATIO^3         | -0.0842          | 0.031               | -2.724 0.007          | -0.145 -0.023    |
| Omnibus:          | 572.978          | Durbin-Watson:      | 0.949                 |                  |
| Prob(Omnibus):    | 0.000            | Jarque-Bera (JB):   | 36189.609             |                  |
| Skew:             | 5.303            | Prob(JB):           | 0.00                  |                  |
| Kurtosis:         | 43.050           | Cond. No.           | 3.02e+06              |                  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.02e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for B:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.144    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.139    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 28.14    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 7.83e-17 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1766.8  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3542.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3558.    |
| <b>Df Model:</b>         | 3                |                            |          |

**Covariance Type:** nonrobust

|              | coef       | std err  | t      | P> t  | [0.025    | 0.975]  |
|--------------|------------|----------|--------|-------|-----------|---------|
| <b>const</b> | 17.9898    | 2.312    | 7.782  | 0.000 | 13.448    | 22.531  |
| <b>B</b>     | -0.0845    | 0.056    | -1.497 | 0.135 | -0.196    | 0.026   |
| <b>B^2</b>   | 0.0002     | 0.000    | 0.760  | 0.447 | -0.000    | 0.001   |
| <b>B^3</b>   | -2.895e-07 | 4.38e-07 | -0.661 | 0.509 | -1.15e-06 | 5.7e-07 |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 589.534 | <b>Durbin-Watson:</b>    | 0.990     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 42752.655 |
| <b>Skew:</b>          | 5.512   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 46.661  | <b>Cond. No.</b>         | 3.59e+08  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.59e+08. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for LSTAT:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.214    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.210    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 45.67    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 4.13e-26 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1745.0  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3498.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3515.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                | coef    | std err | t      | P> t  | [0.025 | 0.975] |
|----------------|---------|---------|--------|-------|--------|--------|
| <b>const</b>   | 1.0836  | 2.032   | 0.533  | 0.594 | -2.909 | 5.076  |
| <b>LSTAT</b>   | -0.4133 | 0.466   | -0.887 | 0.375 | -1.328 | 0.502  |
| <b>LSTAT^2</b> | 0.0530  | 0.030   | 1.758  | 0.079 | -0.006 | 0.112  |
| <b>LSTAT^3</b> | -0.0008 | 0.001   | -1.423 | 0.155 | -0.002 | 0.000  |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 607.032 | <b>Durbin-Watson:</b>    | 1.239     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 53255.699 |
| <b>Skew:</b>          | 5.717   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 51.941  | <b>Cond. No.</b>         | 5.20e+04  |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for MDEV:

# OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | CRIM             | <b>R-squared:</b>          | 0.416    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.413    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 119.2    |
| <b>Date:</b>             | Wed, 17 Jan 2024 | <b>Prob (F-statistic):</b> | 2.65e-58 |
| <b>Time:</b>             | 14:48:40         | <b>Log-Likelihood:</b>     | -1670.0  |
| <b>No. Observations:</b> | 506              | <b>AIC:</b>                | 3348.    |
| <b>Df Residuals:</b>     | 502              | <b>BIC:</b>                | 3365.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|               | coef    | std err | t       | P> t  | [0.025 | 0.975] |
|---------------|---------|---------|---------|-------|--------|--------|
| <b>const</b>  | 52.9386 | 3.366   | 15.725  | 0.000 | 46.325 | 59.553 |
| <b>MDEV</b>   | -5.0774 | 0.435   | -11.668 | 0.000 | -5.932 | -4.222 |
| <b>MDEV^2</b> | 0.1551  | 0.017   | 8.995   | 0.000 | 0.121  | 0.189  |
| <b>MDEV^3</b> | -0.0015 | 0.000   | -7.277  | 0.000 | -0.002 | -0.001 |

|                       |         |                          |           |
|-----------------------|---------|--------------------------|-----------|
| <b>Omnibus:</b>       | 568.100 | <b>Durbin-Watson:</b>    | 1.360     |
| <b>Prob(Omnibus):</b> | 0.000   | <b>Jarque-Bera (JB):</b> | 47296.533 |
| <b>Skew:</b>          | 5.084   | <b>Prob(JB):</b>         | 0.00      |
| <b>Kurtosis:</b>      | 49.259  | <b>Cond. No.</b>         | 3.67e+05  |

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.67e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Several predictors have statistically significant polynomial terms at the 5% level, including the proportion of indutrially zoned acres, nitric oxide concentration, proportion of homes built before 1940, distance from employment centers, pupil-teacher ratio, and median home value. This makes sense if we look at the plots from part (a) since many of the relationships are flat and then spike up. These relationships can be modeled better with a curve than with a straight line, though curves are still not perfect. A model that can spike up and down like KNN might be able to fit the data better. We could also try taking the natural log of the per capita crime rate to reduce the spikiness.



