

PPHA 30546 Machine Learning Problem Set 1

Richard Campo

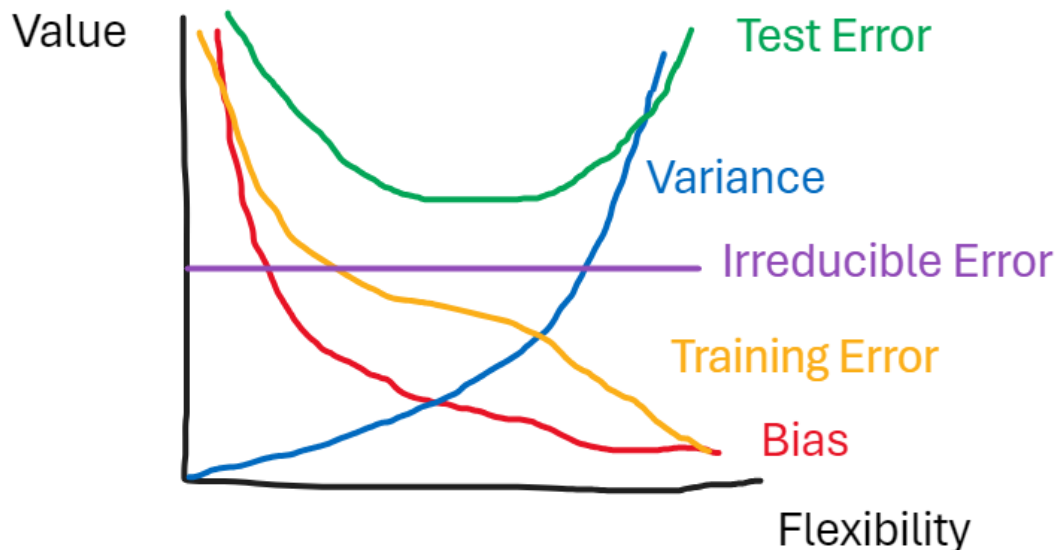
```
In [ ]: from IPython.display import Image
import pandas as pd
import os
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
from IPython.display import display
```

Chapter 2: Question 3

(a)

```
In [37]: # https://stackoverflow.com/a/11855133
Image(filename="./Machine Learning Pset 1 Curves.png")
```

Out[37]:



(b)

- Bias decreases as flexibility increases because the model can overfit the data. Bias tends to decrease faster at lower levels of flexibility because adding just a little curvature can allow the model to fit the data much better, but once you have lots of flexibility, adding more won't do much.
- Variance increases as flexibility increases because the model is more sensitive to new training data. Variance tends to increase at an increasing rate because very flexible models will change more than less

flexible models when provided new training data due to overfitting.

- Training error decreases as flexibility increases because the model can overfit the data. Eventually, the most flexible model can correctly predict every point in the data by overfitting.
- Test error first decreases and then increases as flexibility increases. When flexibility is low, adding flexibility allows the model to more closely fit the training data and accurately predict the test data. However, when flexibility is high, adding flexibility causes overfitting, which lowers training error but increases test error.
- Irreducible error is a constant = $\text{Var}(\epsilon)$

Chapter 2: Question 5

The advantage of a very flexible approach is that the model can more closely fit the training data and more accurately predict the test data. However, the disadvantage of a very flexible approach is that the model can overfit the training data, meaning the model will make very accurate predictions on the training data, but not generalize to the test data.

A more flexible approach is preferred when the true underlying relationship in the data is very nonlinear because a line can't fit the data well. On the other hand, a less flexible approach is preferred when the true underlying relationship in the data is linear or approximately linear because flexible methods will overfit the data.

Chapter 2: Question 10

(a)

```
In [38]: PATH = r"C:\Users\RichardCampo\Documents\GitHub\Machine-Learning\Boston"
df_boston = pd.read_csv(os.path.join(PATH, "Boston.csv"))
```

(b)

```
In [39]: print(df_boston.shape)
df_boston.head()
```

(506, 14)

```
Out[39]:
```

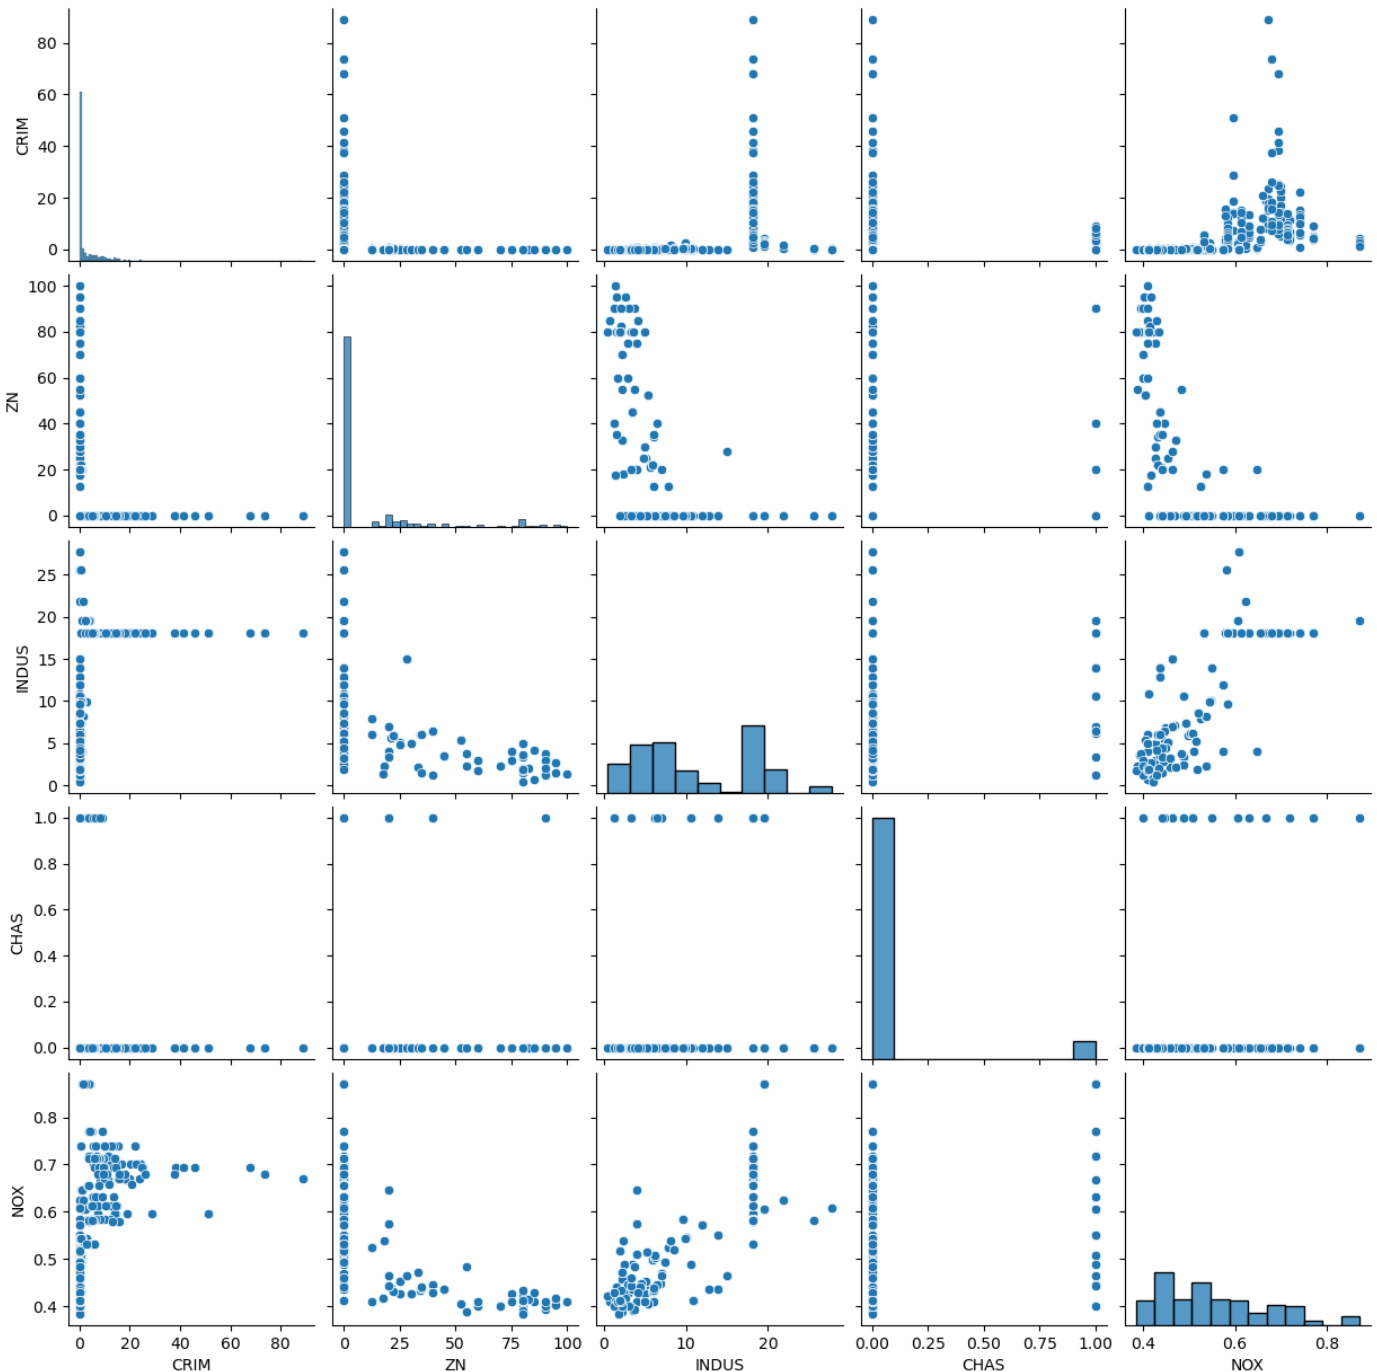
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MDEV
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

There are 506 rows and 14 columns in the dataset. Rows represent census tracts in Boston and columns represent variables describing each census tract.

(c)

```
In [40]: sns.pairplot(df_boston.iloc[:, :5])

Out[40]: <seaborn.axisgrid.PairGrid at 0x21e75e695d0>
```



Looking at the scatterplots, there is a positive correlations between nitric oxide concentration and the number of acres zoned for industry in a census district, which is unsurprising. It also looks like there is a negative association between the number of acres zoned for residential buildings and the amount of nitric oxide and amount of acres zoned for industry. This also makes sense. Lastly, per capita crime appears to be positively correlated with the concentration of nitric oxides.

(d)

```
In [41]: df_boston.corr()["CRIM"].sort_values(ascending=False)
```

```
Out[41]: CRIM    1.000000
RAD      0.622029
TAX      0.579564
LSTAT    0.452220
NOX      0.417521
```

```

INDUS      0.404471
AGE        0.350784
PTRATIO    0.288250
CHAS       -0.055295
ZN         -0.199458
RM         -0.219940
B          -0.377365
DIS        -0.377904
MDEV       -0.385832
Name: CRIM, dtype: float64

```

Most of the predictors are at least weakly associated with crime. The strongest correlation is between per capita crime and access to radial highways, maybe because areas close to highways are poorer or maybe because criminals use highways to move illegal drugs. Property tax rate is also positively associated with crime, perhaps because property taxes are higher closer to the city center where crime is more prevalent. Lastly, LSTAT, the percentage of people in the census district who are "lower status" (presumably in terms of income), is also positively correlated with per capita crime, which is what we expect.

(e)

```

In [42]: print(df_boston[["CRIM", "TAX", "PTRATIO"]].describe())
df_boston[["CRIM", "TAX", "PTRATIO"]].hist()

```

```

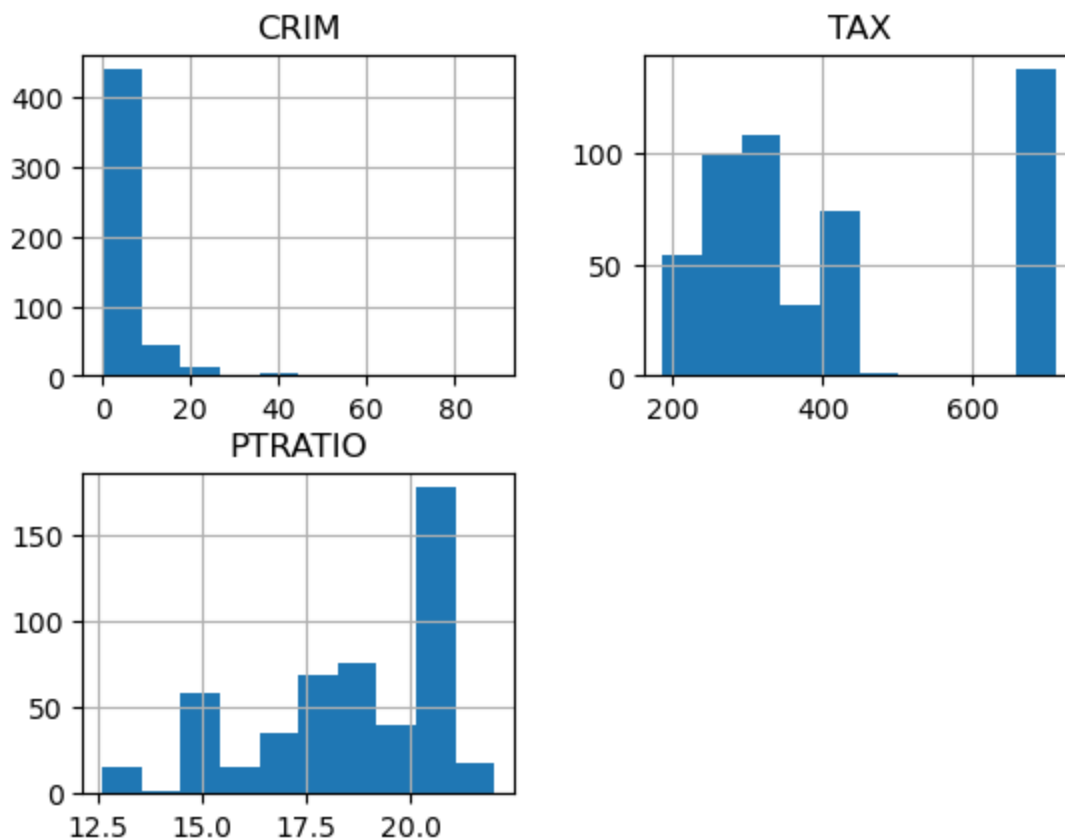
          CRIM          TAX          PTRATIO
count  506.000000  506.000000  506.000000
mean     3.593761  408.237154   18.455534
std     8.596783  168.537116    2.164946
min     0.006320  187.000000   12.600000
25%     0.082045  279.000000   17.400000
50%     0.256510  330.000000   19.050000
75%     3.647423  666.000000   20.200000
max     88.976200  711.000000   22.000000

```

```

Out[42]: array([[<Axes: title={'center': 'CRIM'}>,
      <Axes: title={'center': 'TAX'}>],
      [<Axes: title={'center': 'PTRATIO'}>, <Axes: >]], dtype=object)

```



Most census tracts have low crime per capita since the mean is only about 3.6. However, a small number of census tracts have high per capita crime rates, up to a maximum of almost 89.

Census tracts in Boston have full-value property tax rates per \$10,000 that vary from 187 to 711, but there is a large valley between about 450 and 650. A large group of census tracts have high property tax rates, possibly because they are close to the city center.

Lastly, pupil-teacher ratios vary from 12.6 to 22 pupils per teacher, with a large number of census districts having 20 to 21 pupils per teacher. These higher ratio census districts may also be closer to the city center where schools tend to have less funding per student.

(f)

```
In [43]: borders = len(df_boston[df_boston["CHAS"] == 1])

print(borders)
print(borders/len(df_boston) * 100, "%")
```

```
35
6.91699604743083 %
```

35 of the 506 census tracts border the Charles River, or about 7% of census tracts in the dataset.

(g)

```
In [44]: df_boston["PTRATIO"].median()
```

```
Out[44]: 19.05
```

The median pupil-teacher ratio of census tracts in the dataset is 19.05 pupils per teacher.

(h)

```
In [45]: df_boston.loc[[df_boston["MDEV"].idxmin()]]
```

```
Out[45]:
```

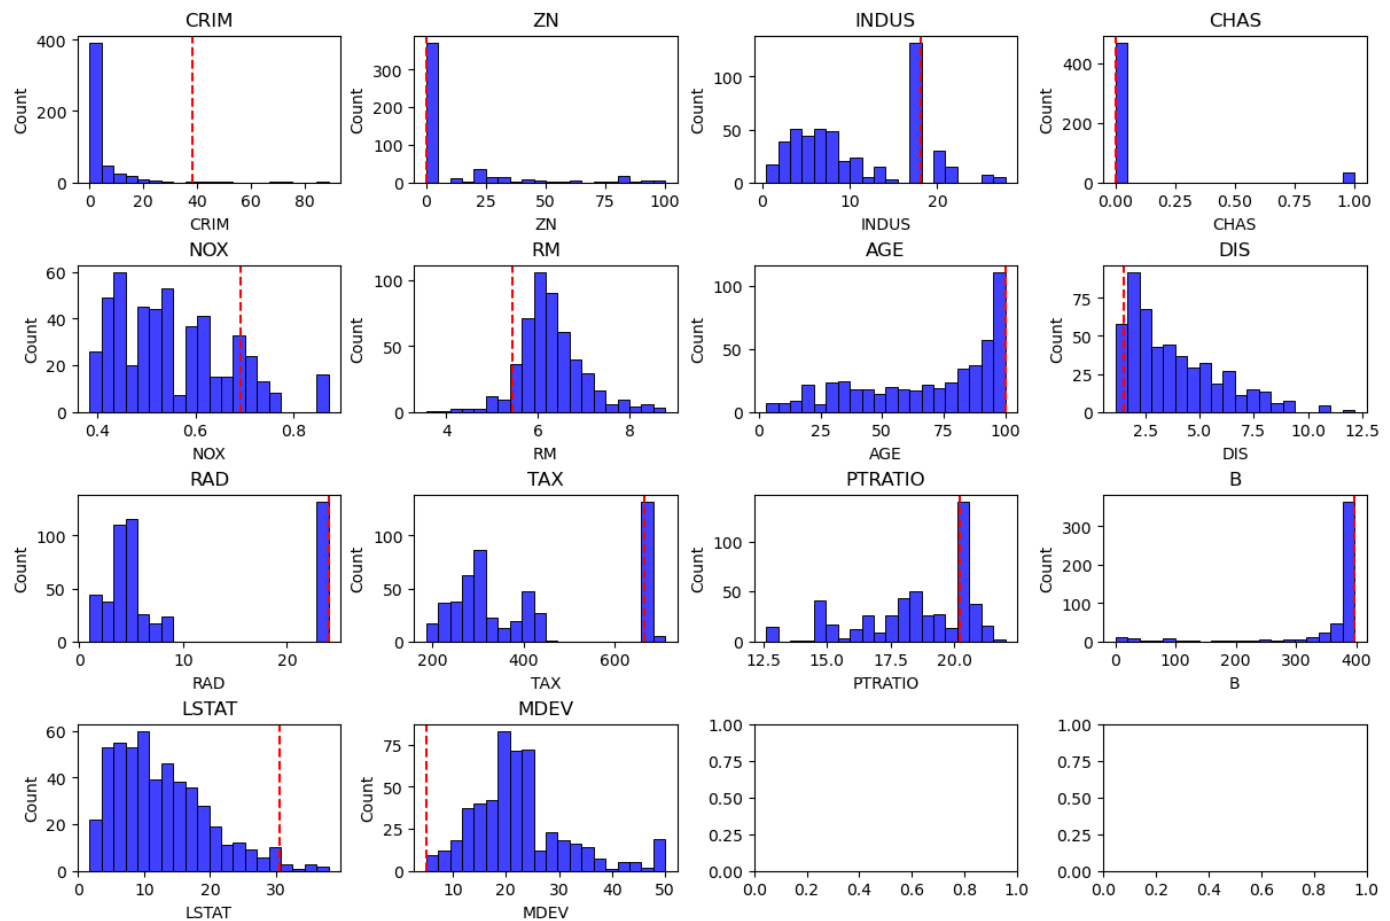
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MDEV
398	38.3518	0.0	18.1	0.0	0.693	5.453	100.0	1.4896	24.0	666.0	20.2	396.9	30.59	5.0

```
In [46]: min_row_index = df_boston["MDEV"].idxmin()
min_row = df_boston.loc[[min_row_index]]

# https://stackoverflow.com/a/53649492
fig, axs = plt.subplots(4, 4, figsize=(12, 8), constrained_layout=True)
fig.suptitle("Histograms with Min MDEV Tract Highlighted", fontsize=16, y=1.05)

# Plot a histogram showing distribution for each variable with red line
# showing the value for the census tract with lowest MDEV.
for ax, column in zip(axs.flatten(), df_boston.columns):
    sns.histplot(df_boston[column], bins=20, color='blue', ax=ax)
    ax.axvline(x=min_row[column].values[0], color='red', linestyle='dashed')
    ax.set_title(column)
```

Histograms with Min MDEV Tract Highlighted



The census tract with the lowest median value of owner-occupied homes is located in row 398 and has a median home value of 5,000. *This is much lower than the median value of 21.2 thousand.* This census tract also has high crime per capita, an above average amount of industrially zoned space, above average nitric oxide concentration, low average number of rooms per building, very old buildings, low distance to the city employment centers, high accessibility to radial highways, high property taxes, a high pupil-teacher ratio, and a high proportion of people with low incomes.

(i)

```
In [47]: print(len(df_boston.loc[df_boston["RM"] > 7]))
print(len(df_boston.loc[df_boston["RM"] > 8]))
```

```
64
13
```

```
In [48]: more_than_8 = df_boston.loc[df_boston["RM"] > 8].mean()
avg_values = df_boston.mean()

comparison = pd.concat([more_than_8, avg_values], axis = 1)
comparison.columns = ["more_than_8", "avg_values"]

print(comparison)
```

	more_than_8	avg_values
CRIM	0.718795	3.593761
ZN	13.615385	11.363636
INDUS	7.078462	11.136779
CHAS	0.153846	0.069170
NOX	0.539238	0.554695
RM	8.348538	6.284634
AGE	71.538462	68.574901

DIS	3.430192	3.795043
RAD	7.461538	9.549407
TAX	325.076923	408.237154
PTRATIO	16.361538	18.455534
B	385.210769	356.674032
LSTAT	4.310000	12.653063
MDEV	44.200000	22.532806

Census tracts with more than eight rooms per dwelling tend to have low crime, have less industrially zoned land, are more likely to border the Charles River, have a smaller share of low income residents, and have high median home values.

Chapter 3: Question 3

(a)

The "true" regression model is:

$$salary = \beta_0 + \beta_1 GPA + \beta_2 IQ + \beta_3 level + \beta_4 GPA \times IQ + \beta_5 GPA \times level + \epsilon$$

We estimate the model:

$$\widehat{salary} = 50 + 20(GPA) + 0.07(IQ) + 35(level) + 0.01(GPA \times IQ) - 10(GPA \times level) + e$$

The answer is (ii). For a fixed value of IQ and GPA, college graduates earn more on average than high school graduates. This is because β_3 , the coefficient of level is 35, meaning a college graduate can expect a \$35,000 higher starting salary than a high school graduate on average, all else constant. The estimated coefficient of -10 for the interaction of GPA and level indicates that for college graduates, the slope of GPA is more shallow, but we are interested in the average difference between levels for this question, not GPA.

(b)

Based on our model, we predict that the starting salary of a college graduate with an IQ of 110 and a GPA of 4.0 will be

$$50 + 20(4.0) + 0.07(110) + 35(1) + 0.01(4.0)(110) - 10(4.0)(1) = 137.1$$

or \$137.1 thousand.

(c)

False. Just because the coefficient for the interaction between GPA and IQ is small does not mean there is little evidence of an interaction effect. The formula for a t-test to check if the coefficient is statistically significant is:

$$t = \frac{\hat{\beta}_j - \beta_j}{SE(\hat{\beta}_j)}$$

This means that even if the estimated coefficient is small, as long as the standard error is also very small, the coefficient could be statistically significant.

Chapter 3: Question 15

(a)

```
In [49]: predictors = list(df_boston.columns)
predictors.remove("CRIM")

univar_coefs = {} # for part (c)

for predictor in predictors:
    X = df_boston[predictor]
    X = sm.add_constant(X) # add intercept to regression model
    y = df_boston["CRIM"]

    model = sm.OLS(y, X).fit()

    # Extract coefficient variable names and values for part (c)
    univar_coefs[predictor] = model.params[predictor]

print(f"Regression results for {predictor}:")
display(model.summary())
print("\n")
```

Regression results for ZN:

OLS Regression Results						
Dep. Variable:		CRIM		R-squared:		0.040
Model:		OLS		Adj. R-squared:		0.038
Method:		Least Squares		F-statistic:		20.88
Date:		Sat, 13 Jan 2024		Prob (F-statistic):		6.15e-06
Time:		13:58:23		Log-Likelihood:		-1795.8
No. Observations:		506		AIC:		3596.
Df Residuals:		504		BIC:		3604.
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]
const	4.4292	0.417	10.620	0.000	3.610	5.249
ZN	-0.0735	0.016	-4.570	0.000	-0.105	-0.042
Omnibus:		568.366		Durbin-Watson:		0.862
Prob(Omnibus):		0.000		Jarque-Bera (JB):		32952.356
Skew:		5.270		Prob(JB):		0.00
Kurtosis:		41.103		Cond. No.		28.8

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for INDUS:

OLS Regression Results				
Dep. Variable:	CRIM		R-squared:	0.164
Model:	OLS		Adj. R-squared:	0.162

Method:	Least Squares	F-statistic:	98.58
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	2.44e-21
Time:	13:58:23	Log-Likelihood:	-1760.9
No. Observations:	506	AIC:	3526.
Df Residuals:	504	BIC:	3534.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-2.0509	0.668	-3.072	0.002	-3.362	-0.739
INDUS	0.5068	0.051	9.929	0.000	0.407	0.607

Omnibus:	585.528	Durbin-Watson:	0.990
Prob(Omnibus):	0.000	Jarque-Bera (JB):	41469.710
Skew:	5.456	Prob(JB):	0.00
Kurtosis:	45.987	Cond. No.	25.1

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for CHAS:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.003
Model:	OLS	Adj. R-squared:	0.001
Method:	Least Squares	F-statistic:	1.546
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	0.214
Time:	13:58:23	Log-Likelihood:	-1805.3
No. Observations:	506	AIC:	3615.
Df Residuals:	504	BIC:	3623.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	3.7232	0.396	9.404	0.000	2.945	4.501
CHAS	-1.8715	1.505	-1.243	0.214	-4.829	1.086

Omnibus:	562.698	Durbin-Watson:	0.822
Prob(Omnibus):	0.000	Jarque-Bera (JB):	30864.755
Skew:	5.205	Prob(JB):	0.00
Kurtosis:	39.818	Cond. No.	3.96

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for NOX:

OLS Regression Results						
Dep. Variable:		CRIM		R-squared:		0.174
Model:		OLS		Adj. R-squared:		0.173
Method:		Least Squares		F-statistic:		106.4
Date:		Sat, 13 Jan 2024		Prob (F-statistic):		9.16e-23
Time:		13:58:23		Log-Likelihood:		-1757.6
No. Observations:		506		AIC:		3519.
Df Residuals:		504		BIC:		3528.
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]
const	-13.5881	1.702	-7.986	0.000	-16.931	-10.245
NOX	30.9753	3.003	10.315	0.000	25.076	36.875
Omnibus:		591.496		Durbin-Watson:		0.994
Prob(Omnibus):		0.000		Jarque-Bera (JB):		42994.381
Skew:		5.544		Prob(JB):		0.00
Kurtosis:		46.776		Cond. No.		11.3

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for RM:

OLS Regression Results						
Dep. Variable:		CRIM		R-squared:		0.048
Model:		OLS		Adj. R-squared:		0.046
Method:		Least Squares		F-statistic:		25.62
Date:		Sat, 13 Jan 2024		Prob (F-statistic):		5.84e-07
Time:		13:58:23		Log-Likelihood:		-1793.5
No. Observations:		506		AIC:		3591.
Df Residuals:		504		BIC:		3600.
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]

const	20.5060	3.362	6.099	0.000	13.901	27.111
RM	-2.6910	0.532	-5.062	0.000	-3.736	-1.646
Omnibus:	576.890	Durbin-Watson:	0.883			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36966.825			
Skew:	5.361	Prob(JB):	0.00			
Kurtosis:	43.477	Cond. No.	58.4			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for AGE:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.123
Model:	OLS	Adj. R-squared:	0.121
Method:	Least Squares	F-statistic:	70.72
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	4.26e-16
Time:	13:58:23	Log-Likelihood:	-1772.9
No. Observations:	506	AIC:	3550.
Df Residuals:	504	BIC:	3558.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-3.7527	0.944	-3.974	0.000	-5.608	-1.898
AGE	0.1071	0.013	8.409	0.000	0.082	0.132
Omnibus:	575.090	Durbin-Watson:	0.960			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36851.412			
Skew:	5.331	Prob(JB):	0.00			
Kurtosis:	43.426	Cond. No.	195.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for DIS:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.143
Model:	OLS	Adj. R-squared:	0.141
Method:	Least Squares	F-statistic:	83.97

Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.27e-18
Time:	13:58:23	Log-Likelihood:	-1767.1
No. Observations:	506	AIC:	3538.
Df Residuals:	504	BIC:	3547.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	9.4489	0.731	12.934	0.000	8.014	10.884
DIS	-1.5428	0.168	-9.163	0.000	-1.874	-1.212
Omnibus:	577.090	Durbin-Watson:	0.957			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	37542.100			
Skew:	5.357	Prob(JB):	0.00			
Kurtosis:	43.815	Cond. No.	9.32			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for RAD:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.387
Model:	OLS	Adj. R-squared:	0.386
Method:	Least Squares	F-statistic:	318.1
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.62e-55
Time:	13:58:23	Log-Likelihood:	-1682.3
No. Observations:	506	AIC:	3369.
Df Residuals:	504	BIC:	3377.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-2.2709	0.445	-5.105	0.000	-3.145	-1.397
RAD	0.6141	0.034	17.835	0.000	0.546	0.682
Omnibus:	654.232	Durbin-Watson:	1.336			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	74327.568			
Skew:	6.441	Prob(JB):	0.00			
Kurtosis:	60.961	Cond. No.	19.2			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for TAX:

OLS Regression Results						
Dep. Variable:		CRIM		R-squared:		0.336
Model:		OLS		Adj. R-squared:		0.335
Method:		Least Squares		F-statistic:		254.9
Date:		Sat, 13 Jan 2024		Prob (F-statistic):		9.76e-47
Time:		13:58:23		Log-Likelihood:		-1702.5
No. Observations:		506		AIC:		3409.
Df Residuals:		504		BIC:		3418.
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]
const	-8.4748	0.818	-10.365	0.000	-10.081	-6.868
TAX	0.0296	0.002	15.966	0.000	0.026	0.033
Omnibus:		634.003		Durbin-Watson:		1.252
Prob(Omnibus):		0.000		Jarque-Bera (JB):		63141.063
Skew:		6.134		Prob(JB):		0.00
Kurtosis:		56.332		Cond. No.		1.16e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.16e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for PTRATIO:

OLS Regression Results			
Dep. Variable:	CRIM	R-squared:	0.083
Model:	OLS	Adj. R-squared:	0.081
Method:	Least Squares	F-statistic:	45.67
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	3.88e-11
Time:	13:58:23	Log-Likelihood:	-1784.1
No. Observations:	506	AIC:	3572.
Df Residuals:	504	BIC:	3581.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-17.5307	3.147	-5.570	0.000	-23.714	-11.347
PTRATIO	1.1446	0.169	6.758	0.000	0.812	1.477
Omnibus:	568.808	Durbin-Watson:	0.909			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	34373.378			
Skew:	5.256	Prob(JB):	0.00			
Kurtosis:	41.985	Cond. No.	160.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for B:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.142
Model:	OLS	Adj. R-squared:	0.141
Method:	Least Squares	F-statistic:	83.69
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.43e-18
Time:	13:58:23	Log-Likelihood:	-1767.2
No. Observations:	506	AIC:	3538.
Df Residuals:	504	BIC:	3547.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	16.2680	1.430	11.376	0.000	13.458	19.078
B	-0.0355	0.004	-9.148	0.000	-0.043	-0.028
Omnibus:	591.626	Durbin-Watson:	1.001			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	43282.465			
Skew:	5.543	Prob(JB):	0.00			
Kurtosis:	46.932	Cond. No.	1.49e+03			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.49e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for LSTAT:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.205
Model:	OLS	Adj. R-squared:	0.203
Method:	Least Squares	F-statistic:	129.6
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	7.12e-27
Time:	13:58:23	Log-Likelihood:	-1748.2
No. Observations:	506	AIC:	3500.
Df Residuals:	504	BIC:	3509.
Df Model:	1		
Covariance Type:	nonrobust		
	coef	std err	t P> t [0.025 0.975]
const	-3.2946	0.695	-4.742 0.000 -4.660 -1.930
LSTAT	0.5444	0.048	11.383 0.000 0.450 0.638
Omnibus:	600.766	Durbin-Watson:	1.184
Prob(Omnibus):	0.000	Jarque-Bera (JB):	49637.173
Skew:	5.638	Prob(JB):	0.00
Kurtosis:	50.193	Cond. No.	29.7

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Regression results for MDEV:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.149
Model:	OLS	Adj. R-squared:	0.147
Method:	Least Squares	F-statistic:	88.15
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	2.08e-19
Time:	13:58:23	Log-Likelihood:	-1765.3
No. Observations:	506	AIC:	3535.
Df Residuals:	504	BIC:	3543.
Df Model:	1		
Covariance Type:	nonrobust		
	coef	std err	t P> t [0.025 0.975]
const	11.7202	0.935	12.539 0.000 9.884 13.557
MDEV	-0.3606	0.038	-9.389 0.000 -0.436 -0.285
Omnibus:	559.282	Durbin-Watson:	1.000
Prob(Omnibus):	0.000	Jarque-Bera (JB):	32809.507
Skew:	5.114	Prob(JB):	0.00

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Almost every predictor on its own is significantly correlated with the per capita crime rate, as we can see from the p-values that are substantially lower than 0.05. The only exception is whether the census tract borders the Charles River, which is weakly correlated with lower per capita crime, but the estimated coefficient is not statistically significant at the 5% level.

```
In [50]: for predictor in predictors:
          X = df_boston[predictor]
          X = sm.add_constant(X)  # add the intercept
          y = df_boston["CRIM"]

          model = sm.OLS(y, X).fit()

          # Create Figure and Axes objects
          fig, ax = plt.subplots(figsize=(6, 4))

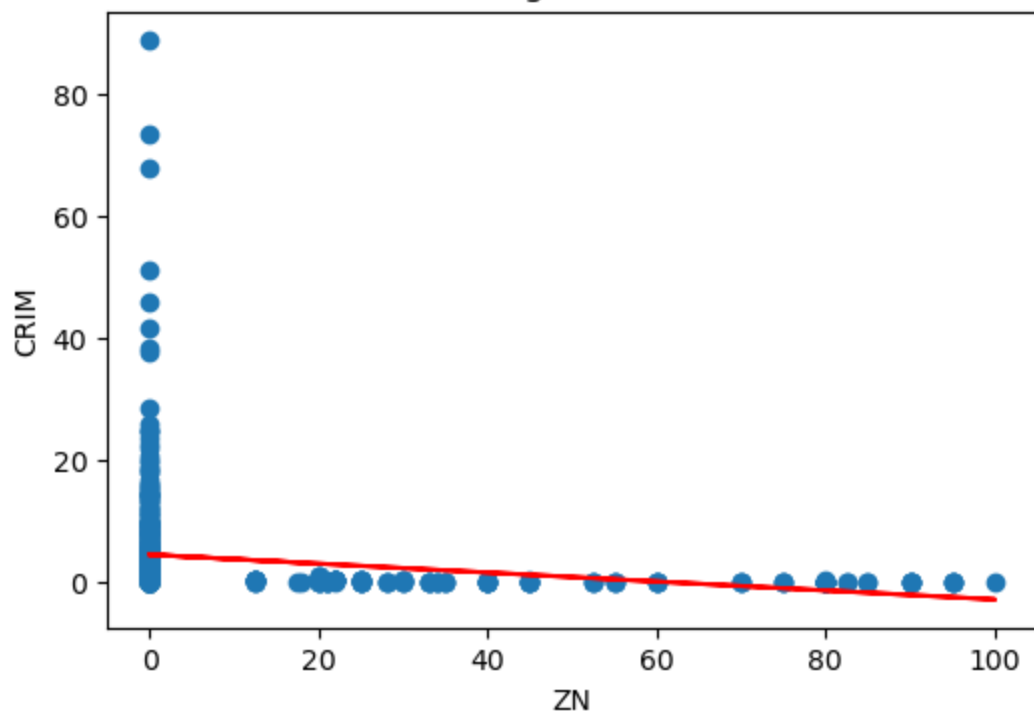
          # Plot the data points
          scatter = ax.scatter(
              X.iloc[:, 1], # exclude intercept column
              y
          )

          # Plot the regression line
          line = ax.plot(
              X.iloc[:, 1], # exclude intercept column
              model.predict(X),
              color='red'
          )

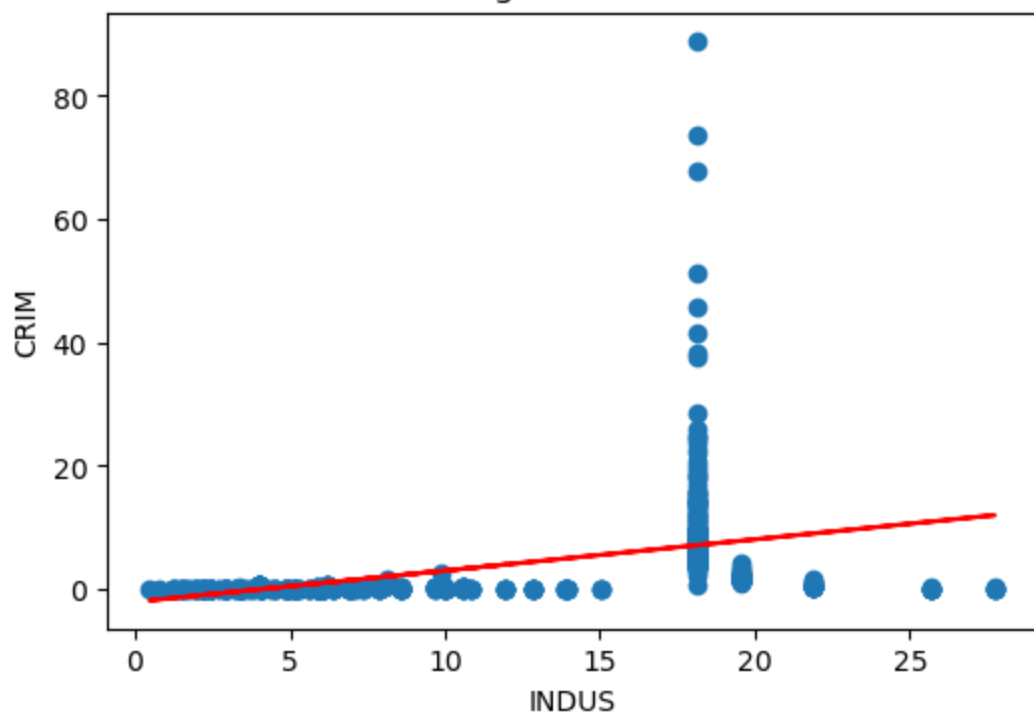
          ax.set_title(f"Linear Regression for {predictor}")
          ax.set_xlabel(predictor)
          ax.set_ylabel("CRIM")

          # Display the plot
          plt.show()
```

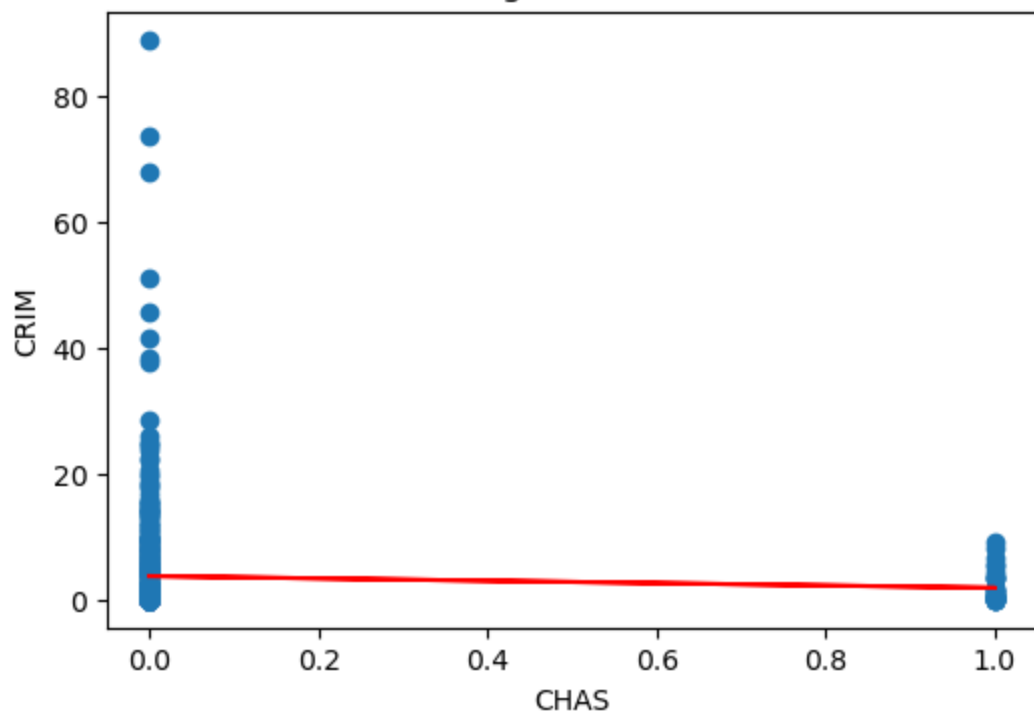

Linear Regression for ZN



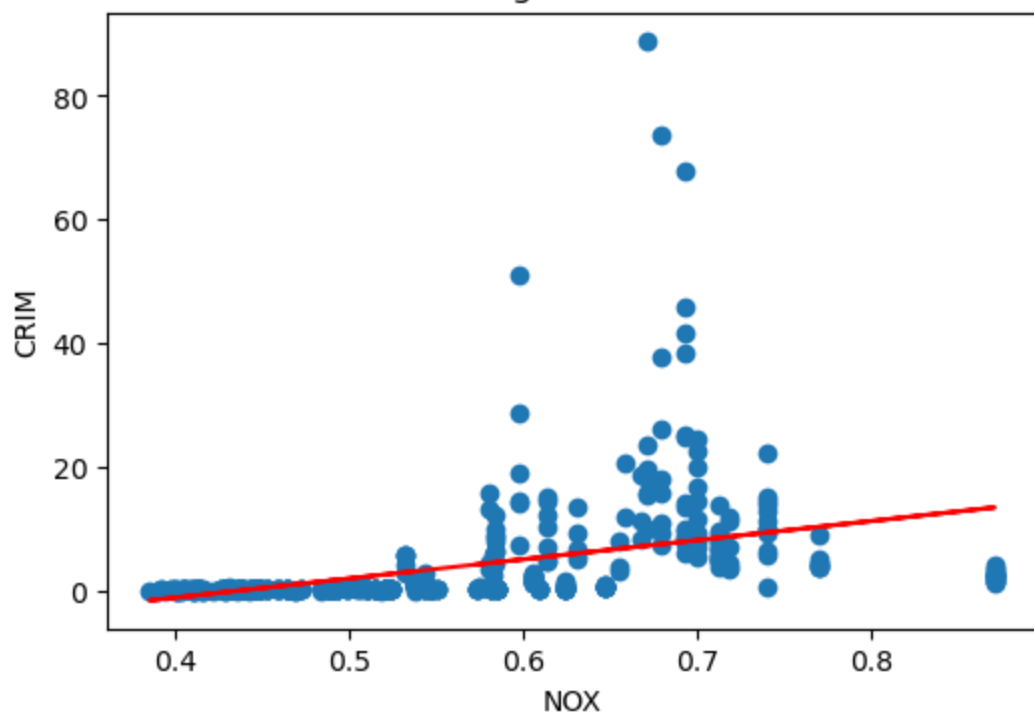
Linear Regression for INDUS



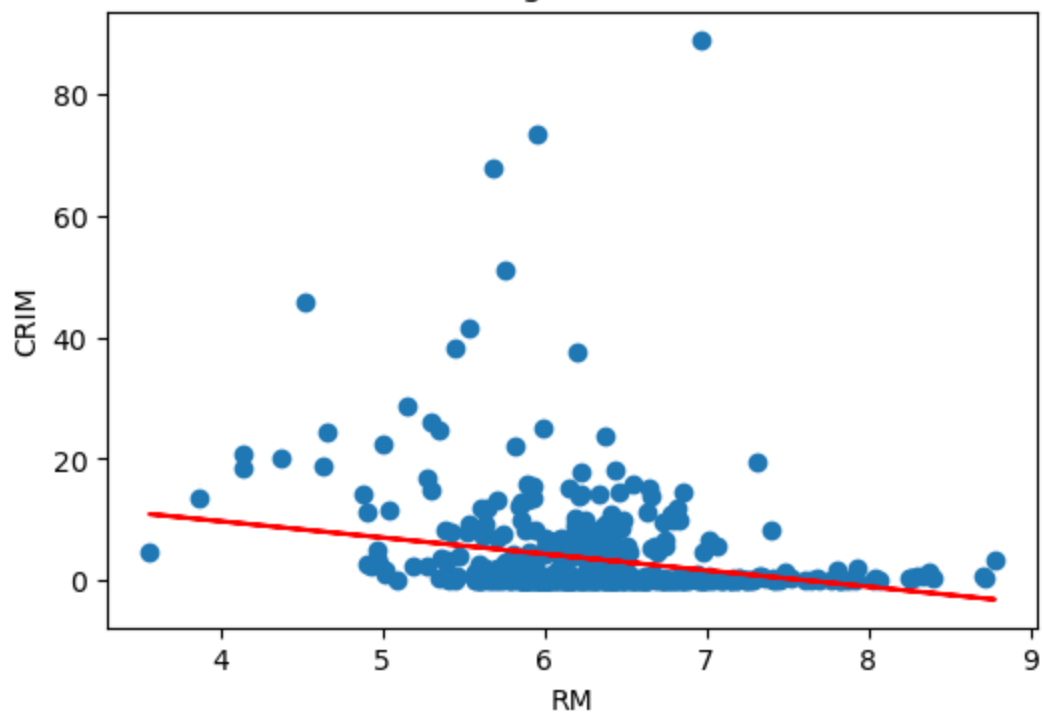
Linear Regression for CHAS



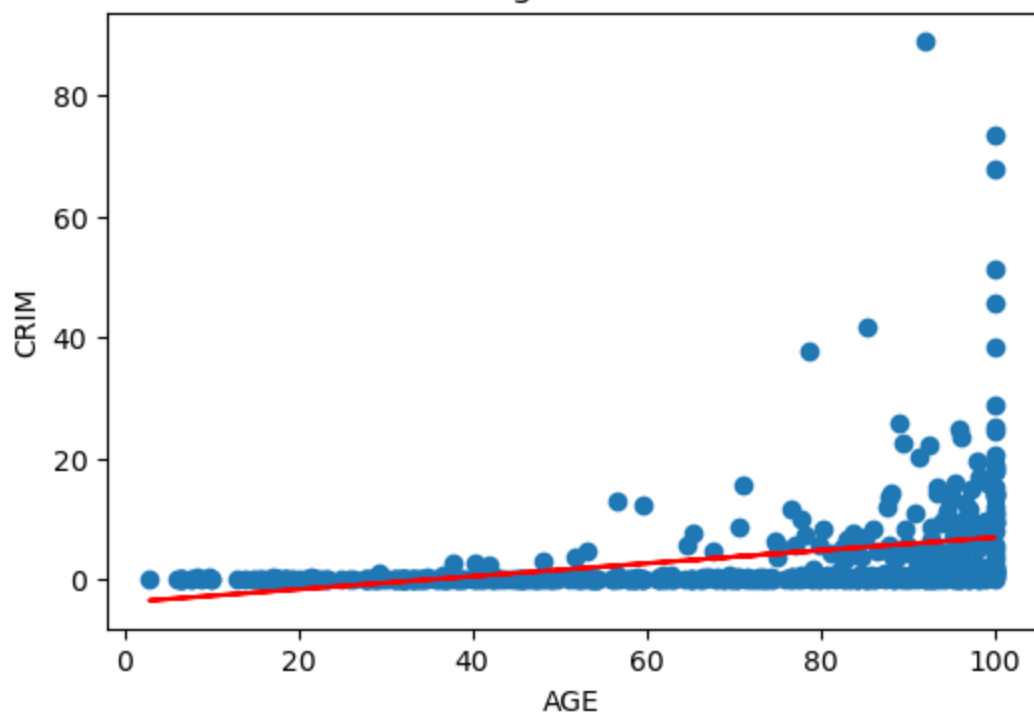
Linear Regression for NOX



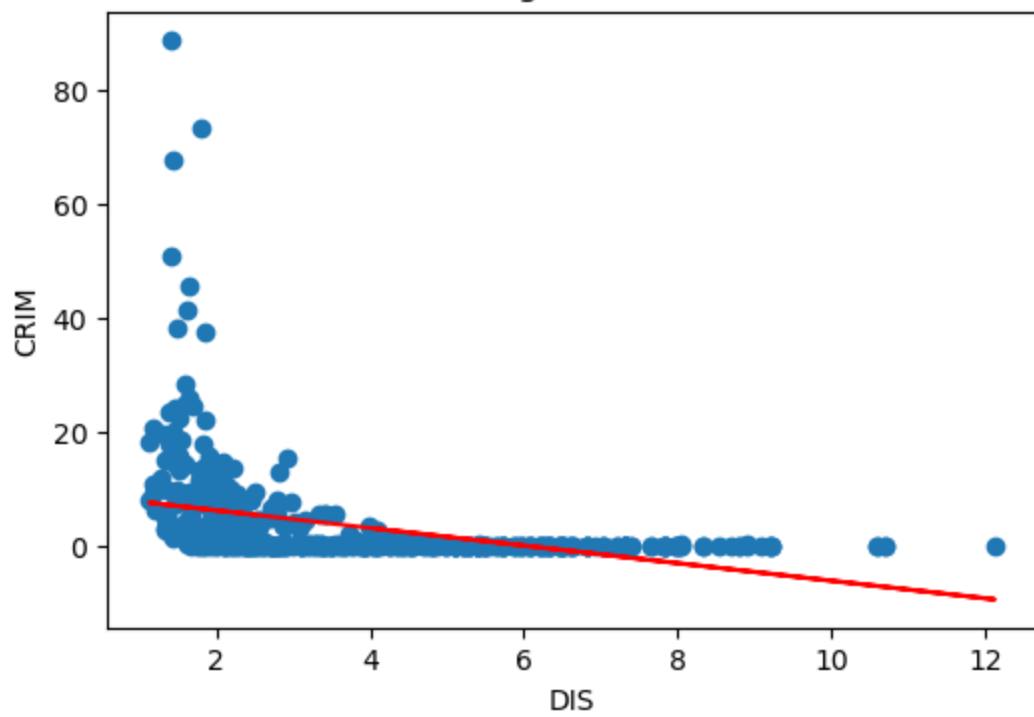
Linear Regression for RM



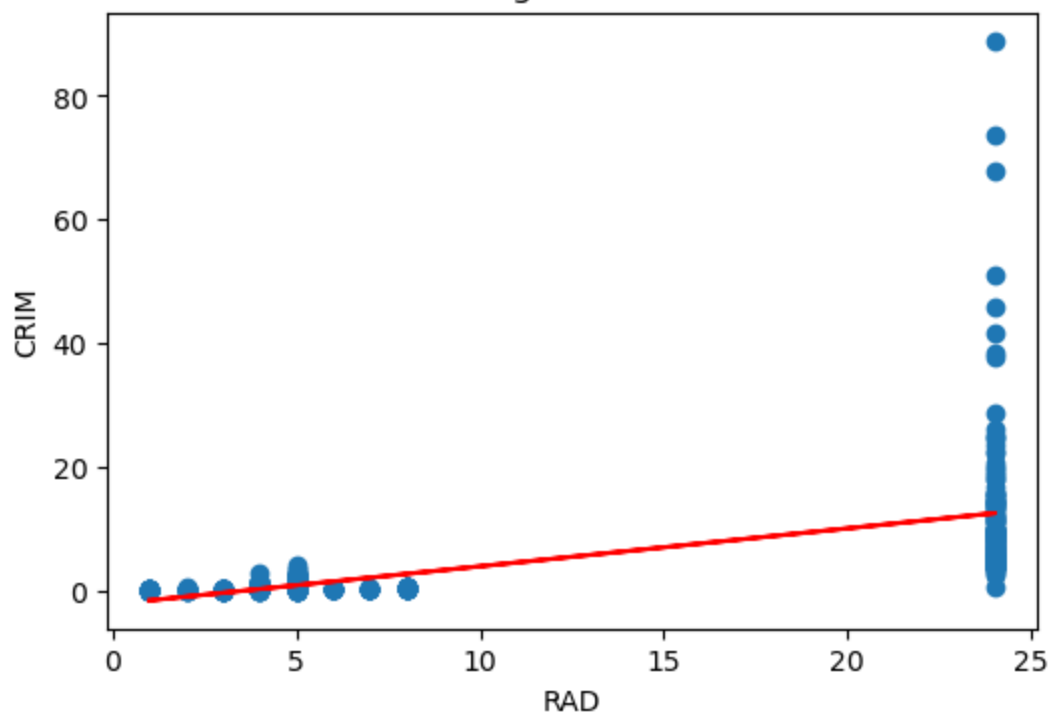
Linear Regression for AGE



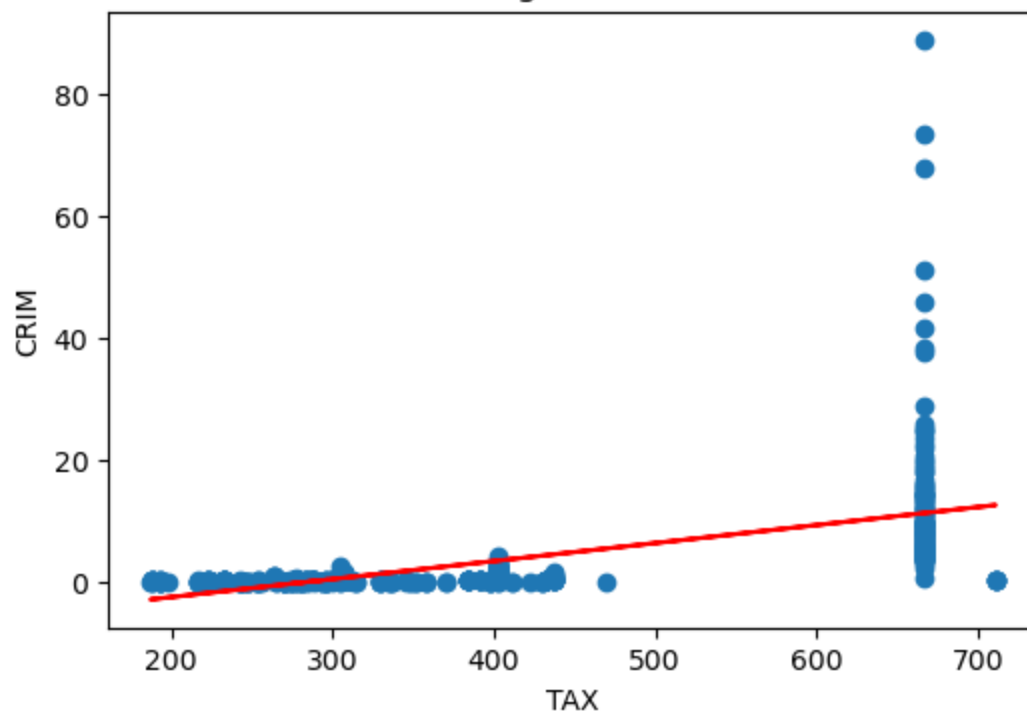
Linear Regression for DIS



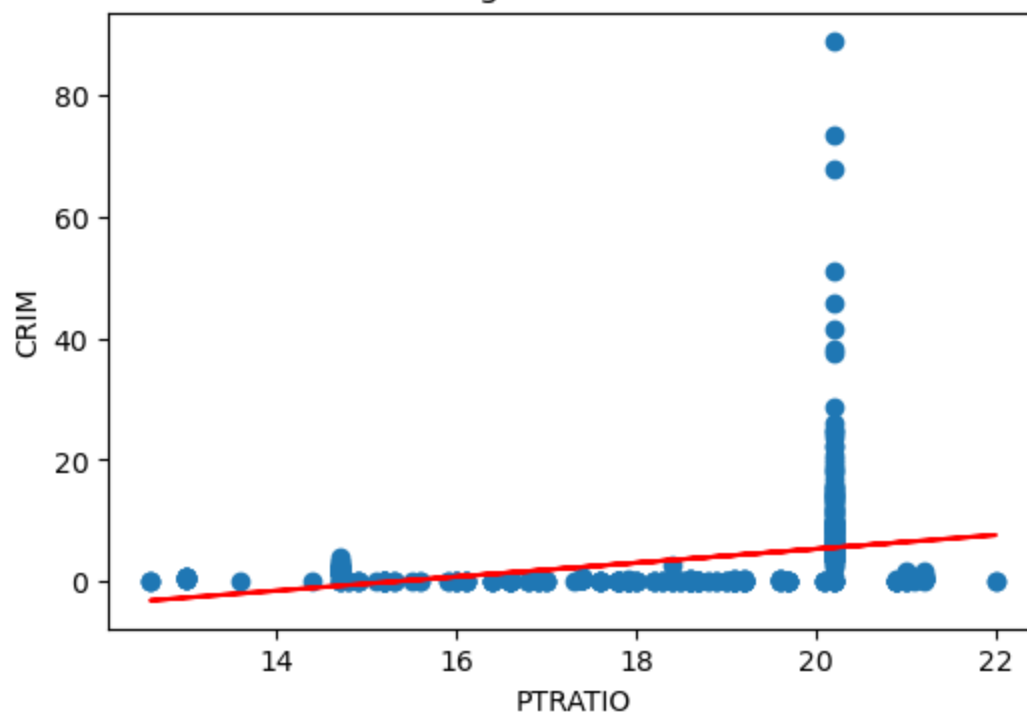
Linear Regression for RAD



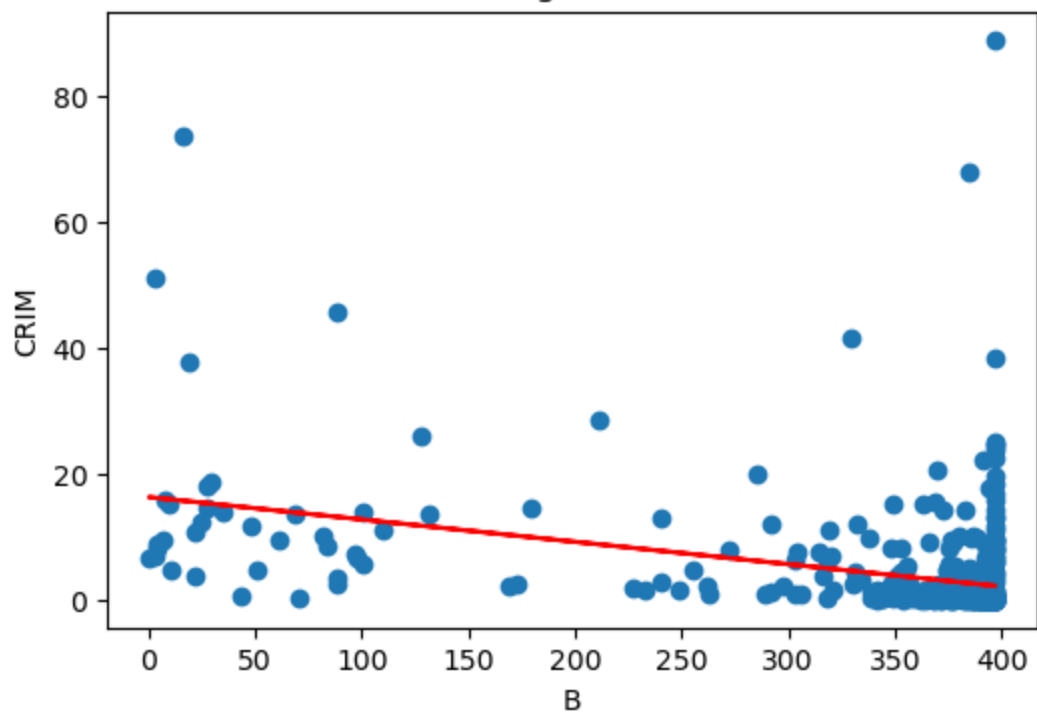
Linear Regression for TAX



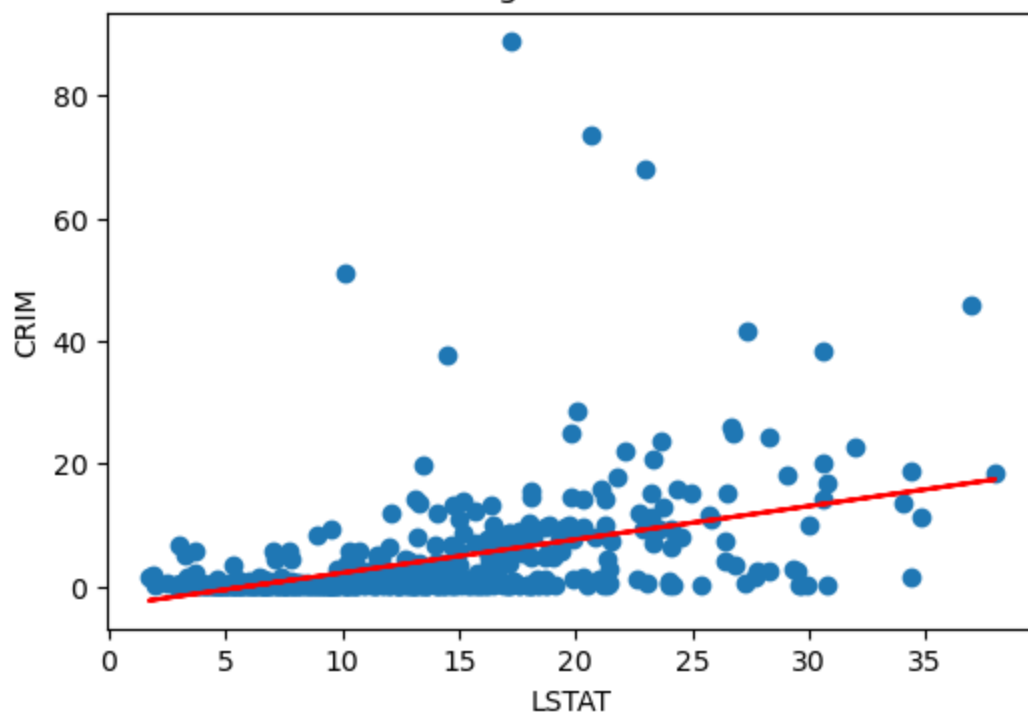
Linear Regression for PTRATIO



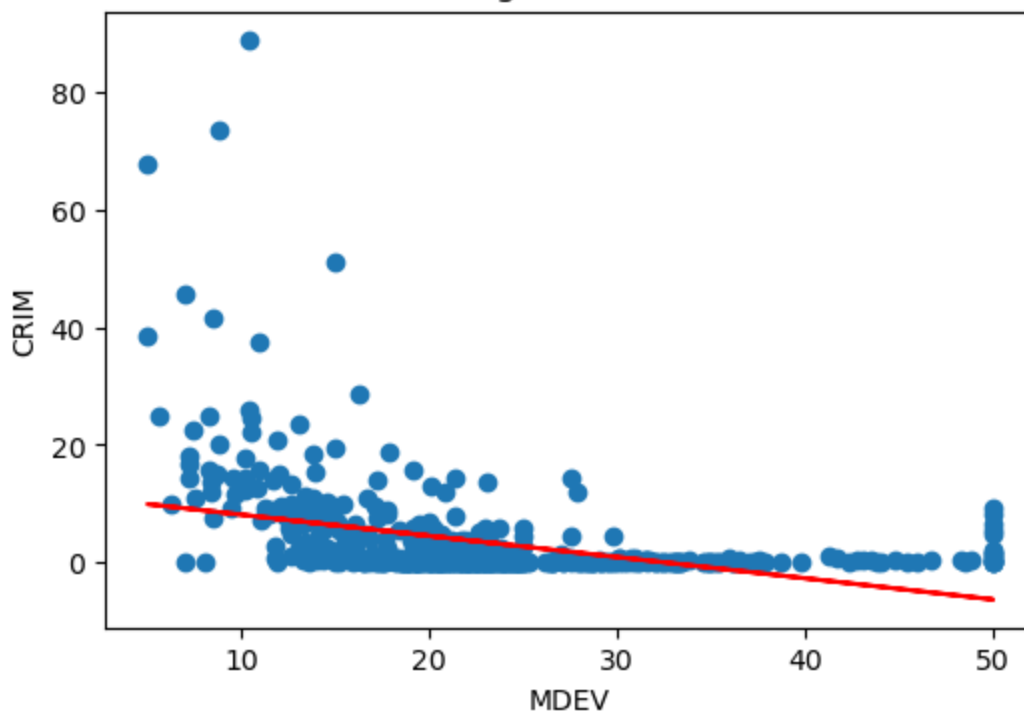
Linear Regression for B



Linear Regression for LSTAT



Linear Regression for MDEV



We can see that there are some correlations for each of the predictors except "CHAS," the dummy for whether the census tract borders the Charles River. There is a weak negative correlation between "CHAS" and "CRIM," but it looks like most census tracts have relatively low crime and do not border the Charles River. The tracts with high crime per capita do not border the Charles River, however.

(b)

```
In [51]: X = df_boston[predictors]
X = sm.add_constant(X) # add intercept to regression model
y = df_boston["CRIM"]

model = sm.OLS(y, X).fit()
model.summary()
```

Out[51]:

OLS Regression Results			
Dep. Variable:	CRIM	R-squared:	0.448
Model:	OLS	Adj. R-squared:	0.434
Method:	Least Squares	F-statistic:	30.73
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	2.04e-55
Time:	13:58:25	Log-Likelihood:	-1655.7
No. Observations:	506	AIC:	3339.
Df Residuals:	492	BIC:	3399.
Df Model:	13		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	17.4184	7.270	2.396	0.017	3.135	31.702
ZN	0.0449	0.019	2.386	0.017	0.008	0.082
INDUS	-0.0616	0.084	-0.735	0.463	-0.226	0.103

CHAS	-0.7414	1.186	-0.625	0.532	-3.071	1.588
NOX	-10.6455	5.301	-2.008	0.045	-21.061	-0.230
RM	0.3811	0.616	0.619	0.536	-0.829	1.591
AGE	0.0020	0.018	0.112	0.911	-0.033	0.037
DIS	-0.9950	0.283	-3.514	0.000	-1.551	-0.439
RAD	0.5888	0.088	6.656	0.000	0.415	0.763
TAX	-0.0037	0.005	-0.723	0.470	-0.014	0.006
PTRATIO	-0.2787	0.187	-1.488	0.137	-0.647	0.089
B	-0.0069	0.004	-1.857	0.064	-0.014	0.000
LSTAT	0.1213	0.076	1.594	0.112	-0.028	0.271
MDEV	-0.1992	0.061	-3.276	0.001	-0.319	-0.080

Omnibus:	662.271	Durbin-Watson:	1.515
Prob(Omnibus):	0.000	Jarque-Bera (JB):	82701.666
Skew:	6.544	Prob(JB):	0.00
Kurtosis:	64.248	Cond. No.	1.58e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.58e+04. This might indicate that there are strong multicollinearity or other numerical problems.

When we just look at the partial effects of each predictor, fewer are predictive of the crime rate per capita. Now, only amount of residential zoning, nitric oxide emissions, distance from employment centers, access to radial highways, and median home value are significant at the 5% significance level.

For these predictors, we can reject the null hypothesis that there is no predicted effect of the independent variable on the per capita crime rate in Boston, holding the other predictors constant. We have evidence in favor of the alternative hypothesis that these predictors do have an expected effect on Boston per capita crime rates, holding the other predictors constant.

```
In [56]: multi_coefs = model.params.drop("const")

# https://stackoverflow.com/a/47561390
uni_vs_multi = pd.DataFrame(
    {"univariate coefficients": pd.Series(univar_coefs),
     "multivariate coefficients": multi_coefs}
)

uni_vs_multi
```

```
Out[56]:
```

	univariate coefficients	multivariate coefficients
ZN	-0.073521	0.044919
INDUS	0.506847	-0.061576
CHAS	-1.871545	-0.741435

NOX	30.975259	-10.645500
RM	-2.691045	0.381070
AGE	0.107131	0.002011
DIS	-1.542831	-0.994992
RAD	0.614137	0.588838
TAX	0.029563	-0.003746
PTRATIO	1.144613	-0.278731
B	-0.035535	-0.006855
LSTAT	0.544406	0.121269
MDEV	-0.360647	-0.199218

```
In [59]: fig, ax = plt.subplots()

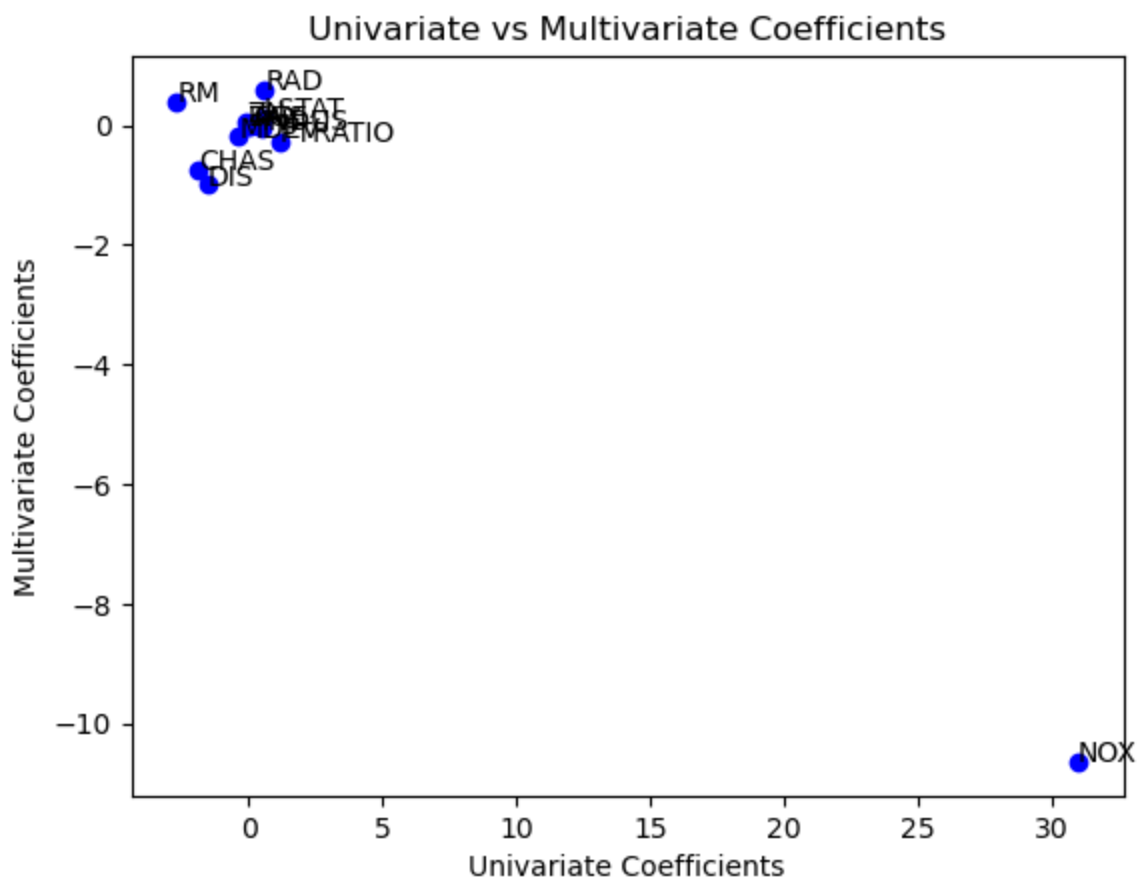
ax.scatter(
    uni_vs_multi['univariate coefficients'],
    uni_vs_multi['multivariate coefficients'],
    color='blue'
)

ax.set_title('Univariate vs Multivariate Coefficients')
ax.set_xlabel('Univariate Coefficients')
ax.set_ylabel('Multivariate Coefficients')

# Add labels for each point
for i, label in enumerate(uni_vs_multi.index):
    ax.annotate(
        label,
        (uni_vs_multi['univariate coefficients'][i],
         uni_vs_multi['multivariate coefficients'][i])
    )

plt.show()

print(df_boston["NOX"].corr(df_boston["RAD"]))
```



0.6114405634855775

We can see from the scatterplot that most of the univariate coefficients were similar though more positive than their multivariate counterparts. The coefficient that changed the most was nitric oxide concentration, which was a strong positive predictor of a higher per capita crime rate in the univariate regression, but actually predicts a lower per capita crime rate when we isolate its partial effect.

This may be because areas with high nitric oxide concentrations tend to be more industrial and closer to the city center, so they tend to have more crime on average. However, high nitric oxide concentrations actually are associated with lower crime rates, all else constant. The difference comes from the fact that NOX is correlated with other variables that are positively associated with crime, such as access to radial highways, which were biasing the univariate model.

(d)

```
In [62]: for predictor in predictors:
          X = df_boston[predictor]
          X_squared = X**2
          X_cubed = X**3

          X_squared = X_squared.rename(f"{predictor}^2")
          X_cubed = X_cubed.rename(f"{predictor}^3")

          X = pd.concat([X, X_squared, X_cubed], axis=1)
          X = sm.add_constant(X)

          y = df_boston["CRIM"]

          model = sm.OLS(y, X).fit()

          print(f"Regression results for {predictor}:")
          display(model.summary())
          print("\n")
```

Regression results for ZN:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.058
Model:	OLS	Adj. R-squared:	0.052
Method:	Least Squares	F-statistic:	10.24
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.49e-06
Time:	14:47:05	Log-Likelihood:	-1791.1
No. Observations:	506	AIC:	3590.
Df Residuals:	502	BIC:	3607.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	4.8193	0.433	11.133	0.000	3.969	5.670
ZN	-0.3303	0.110	-3.008	0.003	-0.546	-0.115
ZN^2	0.0064	0.004	1.670	0.096	-0.001	0.014
ZN^3	-3.753e-05	3.14e-05	-1.196	0.232	-9.92e-05	2.41e-05

Omnibus:	570.003	Durbin-Watson:	0.879
Prob(Omnibus):	0.000	Jarque-Bera (JB):	33886.468
Skew:	5.285	Prob(JB):	0.00
Kurtosis:	41.672	Cond. No.	1.89e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.89e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for INDUS:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.257
Model:	OLS	Adj. R-squared:	0.252
Method:	Least Squares	F-statistic:	57.86
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	3.88e-32
Time:	14:47:05	Log-Likelihood:	-1731.0
No. Observations:	506	AIC:	3470.
Df Residuals:	502	BIC:	3487.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

const	3.6410	1.576	2.310	0.021	0.545	6.737
INDUS	-1.9533	0.483	-4.047	0.000	-2.901	-1.005
INDUS^2	0.2504	0.039	6.361	0.000	0.173	0.328
INDUS^3	-0.0069	0.001	-7.239	0.000	-0.009	-0.005
Omnibus:	611.416	Durbin-Watson:	1.118			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	51547.097			
Skew:	5.815	Prob(JB):	0.00			
Kurtosis:	51.059	Cond. No.	2.47e+04			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.47e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for CHAS:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.003			
Model:	OLS	Adj. R-squared:	-0.001			
Method:	Least Squares	F-statistic:	0.7710			
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	0.463			
Time:	14:47:05	Log-Likelihood:	-1805.3			
No. Observations:	506	AIC:	3617.			
Df Residuals:	503	BIC:	3629.			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	3.7232	0.396	9.395	0.000	2.945	4.502
CHAS	1.108e+14	2.71e+14	0.409	0.683	-4.22e+14	6.43e+14
CHAS^2	-5.578e+13	1.36e+14	-0.409	0.683	-3.24e+14	2.12e+14
CHAS^3	-5.5e+13	1.35e+14	-0.409	0.683	-3.19e+14	2.09e+14
Omnibus:	562.637	Durbin-Watson:	0.822			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	30853.674			
Skew:	5.204	Prob(JB):	0.00			
Kurtosis:	39.811	Cond. No.	5.51e+16			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.7e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Regression results for NOX:

OLS Regression Results			
Dep. Variable:	CRIM	R-squared:	0.292
Model:	OLS	Adj. R-squared:	0.288
Method:	Least Squares	F-statistic:	69.14
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.94e-37
Time:	14:47:05	Log-Likelihood:	-1718.6
No. Observations:	506	AIC:	3445.
Df Residuals:	502	BIC:	3462.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	230.1421	33.734	6.822	0.000	163.864	296.420
NOX	-1264.1021	170.860	-7.398	0.000	-1599.791	-928.414
NOX^2	2223.2265	280.659	7.921	0.000	1671.816	2774.637
NOX^3	-1232.3894	149.687	-8.233	0.000	-1526.479	-938.300

Omnibus:	612.604	Durbin-Watson:	1.159
Prob(Omnibus):	0.000	Jarque-Bera (JB):	52872.508
Skew:	5.824	Prob(JB):	0.00
Kurtosis:	51.705	Cond. No.	1.36e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.36e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for RM:

OLS Regression Results			
Dep. Variable:	CRIM	R-squared:	0.068
Model:	OLS	Adj. R-squared:	0.063
Method:	Least Squares	F-statistic:	12.29
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	9.06e-08
Time:	14:47:05	Log-Likelihood:	-1788.2
No. Observations:	506	AIC:	3584.
Df Residuals:	502	BIC:	3601.
Df Model:	3		

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	111.9002	64.460	1.736	0.083	-14.744	238.545
RM	-38.7040	31.284	-1.237	0.217	-100.167	22.759
RM^2	4.4655	5.005	0.892	0.373	-5.369	14.300
RM^3	-0.1694	0.264	-0.643	0.521	-0.687	0.348

Omnibus:	586.445	Durbin-Watson:	0.919
Prob(Omnibus):	0.000	Jarque-Bera (JB):	40548.719
Skew:	5.484	Prob(JB):	0.00
Kurtosis:	45.461	Cond. No.	5.36e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.36e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for AGE:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.172
Model:	OLS	Adj. R-squared:	0.167
Method:	Least Squares	F-statistic:	34.86
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.76e-20
Time:	14:47:05	Log-Likelihood:	-1758.2
No. Observations:	506	AIC:	3524.
Df Residuals:	502	BIC:	3541.
Df Model:	3		

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-2.5592	2.771	-0.924	0.356	-8.003	2.884
AGE	0.2743	0.186	1.471	0.142	-0.092	0.641
AGE^2	-0.0072	0.004	-1.987	0.047	-0.014	-8.25e-05
AGE^3	5.737e-05	2.11e-05	2.719	0.007	1.59e-05	9.88e-05

Omnibus:	577.859	Durbin-Watson:	1.027
Prob(Omnibus):	0.000	Jarque-Bera (JB):	39629.126
Skew:	5.342	Prob(JB):	0.00
Kurtosis:	45.018	Cond. No.	4.74e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.74e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for DIS:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.276
Model:	OLS	Adj. R-squared:	0.272
Method:	Least Squares	F-statistic:	63.74
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	6.20e-35
Time:	14:47:05	Log-Likelihood:	-1724.4
No. Observations:	506	AIC:	3457.
Df Residuals:	502	BIC:	3474.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	29.9496	2.448	12.235	0.000	25.140	34.759
DIS	-15.5172	1.737	-8.931	0.000	-18.931	-12.104
DIS^2	2.4479	0.347	7.061	0.000	1.767	3.129
DIS^3	-0.1185	0.020	-5.802	0.000	-0.159	-0.078

Omnibus:	577.986	Durbin-Watson:	1.133
Prob(Omnibus):	0.000	Jarque-Bera (JB):	42441.952
Skew:	5.310	Prob(JB):	0.00
Kurtosis:	46.592	Cond. No.	2.10e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.1e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for RAD:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.396
Model:	OLS	Adj. R-squared:	0.392
Method:	Least Squares	F-statistic:	109.5
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	1.47e-54
Time:	14:47:05	Log-Likelihood:	-1678.7

No. Observations:		506		AIC:		3365.	
Df Residuals:		502		BIC:		3382.	
Df Model:		3					
Covariance Type:		nonrobust					
	coef	std err	t	P> t	[0.025	0.975]	
const	-0.6050	2.057	-0.294	0.769	-4.645	3.435	
RAD	0.5122	1.047	0.489	0.625	-1.545	2.569	
RAD^2	-0.0750	0.149	-0.504	0.615	-0.368	0.218	
RAD^3	0.0032	0.005	0.699	0.485	-0.006	0.012	
Omnibus:		657.375		Durbin-Watson:		1.349	
Prob(Omnibus):		0.000		Jarque-Bera (JB):		76643.757	
Skew:		6.487		Prob(JB):		0.00	
Kurtosis:		61.881		Cond. No.		5.43e+04	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.43e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for TAX:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.365		
Model:	OLS	Adj. R-squared:	0.361		
Method:	Least Squares	F-statistic:	96.10		
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	3.69e-49		
Time:	14:47:05	Log-Likelihood:	-1691.3		
No. Observations:	506	AIC:	3391.		
Df Residuals:	502	BIC:	3407.		
Df Model:	3				
Covariance Type:	nonrobust				
	coef	std err	t P> t [0.025 0.975]		
const	19.0705	11.827	1.612 0.107	-4.166 42.307	
TAX	-0.1524	0.096	-1.589 0.113	-0.341 0.036	
TAX^2	0.0004	0.000	1.476 0.141	-0.000 0.001	
TAX^3	-2.193e-07	1.89e-07	-1.158 0.247	-5.91e-07 1.53e-07	
Omnibus:	642.369	Durbin-Watson:	1.292		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	68905.900		

Skew:	6.249	Prob(JB):	0.00
Kurtosis:	58.786	Cond. No.	6.16e+09

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 6.16e+09. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for PTRATIO:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.112
Model:	OLS	Adj. R-squared:	0.107
Method:	Least Squares	F-statistic:	21.21
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	5.99e-13
Time:	14:47:05	Log-Likelihood:	-1775.9
No. Observations:	506	AIC:	3560.
Df Residuals:	502	BIC:	3577.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	474.0255	156.823	3.023	0.003	165.915	782.135
PTRATIO	-81.8089	27.649	-2.959	0.003	-136.131	-27.487
PTRATIO^2	4.6039	1.609	2.862	0.004	1.444	7.764
PTRATIO^3	-0.0842	0.031	-2.724	0.007	-0.145	-0.023

Omnibus:	572.978	Durbin-Watson:	0.949
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36189.609
Skew:	5.303	Prob(JB):	0.00
Kurtosis:	43.050	Cond. No.	3.02e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 3.02e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for B:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.144
Model:	OLS	Adj. R-squared:	0.139
Method:	Least Squares	F-statistic:	28.14

Date:	Sat, 13 Jan 2024	Prob (F-statistic):	7.83e-17
Time:	14:47:05	Log-Likelihood:	-1766.8
No. Observations:	506	AIC:	3542.
Df Residuals:	502	BIC:	3558.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	17.9898	2.312	7.782	0.000	13.448	22.531
B	-0.0845	0.056	-1.497	0.135	-0.196	0.026
B^2	0.0002	0.000	0.760	0.447	-0.000	0.001
B^3	-2.895e-07	4.38e-07	-0.661	0.509	-1.15e-06	5.7e-07

Omnibus:	589.534	Durbin-Watson:	0.990
Prob(Omnibus):	0.000	Jarque-Bera (JB):	42752.655
Skew:	5.512	Prob(JB):	0.00
Kurtosis:	46.661	Cond. No.	3.59e+08

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.59e+08. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for LSTAT:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.214
Model:	OLS	Adj. R-squared:	0.210
Method:	Least Squares	F-statistic:	45.67
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	4.13e-26
Time:	14:47:05	Log-Likelihood:	-1745.0
No. Observations:	506	AIC:	3498.
Df Residuals:	502	BIC:	3515.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	1.0836	2.032	0.533	0.594	-2.909	5.076
LSTAT	-0.4133	0.466	-0.887	0.375	-1.328	0.502
LSTAT^2	0.0530	0.030	1.758	0.079	-0.006	0.112
LSTAT^3	-0.0008	0.001	-1.423	0.155	-0.002	0.000

Omnibus:	607.032	Durbin-Watson:	1.239
Prob(Omnibus):	0.000	Jarque-Bera (JB):	53255.699
Skew:	5.717	Prob(JB):	0.00
Kurtosis:	51.941	Cond. No.	5.20e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Regression results for MDEV:

OLS Regression Results

Dep. Variable:	CRIM	R-squared:	0.416
Model:	OLS	Adj. R-squared:	0.413
Method:	Least Squares	F-statistic:	119.2
Date:	Sat, 13 Jan 2024	Prob (F-statistic):	2.65e-58
Time:	14:47:05	Log-Likelihood:	-1670.0
No. Observations:	506	AIC:	3348.
Df Residuals:	502	BIC:	3365.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	52.9386	3.366	15.725	0.000	46.325	59.553
MDEV	-5.0774	0.435	-11.668	0.000	-5.932	-4.222
MDEV^2	0.1551	0.017	8.995	0.000	0.121	0.189
MDEV^3	-0.0015	0.000	-7.277	0.000	-0.002	-0.001

Omnibus:	568.100	Durbin-Watson:	1.360
Prob(Omnibus):	0.000	Jarque-Bera (JB):	47296.533
Skew:	5.084	Prob(JB):	0.00
Kurtosis:	49.259	Cond. No.	3.67e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.67e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Several predictors have statistically significant polynomial terms at the 5% level, including the proportion of industrial acres, nitric oxide concentration, proportion of homes built before 1940, distance from

employment centers, pupil-teacher ratio, and median home value. This makes sense if we look at the plots from part (a) since many of the relationships are flat and then spike up. These relationships can be modeled better with a curve than with a straight line, though curves are still not perfect. A model that can spike up and down like KNN might be able to fit the data better.